
Probabilistic Programming with Gaussian Process Memoization

Anonymous Author(s)

Affiliation

Address

email

Abstract

This paper describes the *Gaussian process memoizer*, a probabilistic programming technique that uses Gaussian processes to provides a statistical alternative to memorization. Memoizing a target procedure results in a self-caching wrapper that remembers previously computed values. Gaussian process memoization additionally produces a statistical emulator based on a Gaussian process whose predictions automatically improve whenever a new value of the target procedure becomes available. This paper also introduces an efficient implementation, named `gpmem`, that can use kernels given by a broad class of probabilistic programs. The flexibility of `gpmem` is illustrated via three applications: (i) GP regression with hierarchical hyper-parameter learning, (ii) Bayesian structure learning via compositional kernels generated by a probabilistic grammar, and (iii) a bandit formulation of Bayesian optimization with automatic inference and action selection. All applications share a single 50-line Python library and require fewer than 20 lines of probabilistic code each.

1 Introduction

Probabilistic programming could be revolutionary for machine intelligence due to universal inference engines and the rapid prototyping for novel models (Ghahramani, 2015). This levitates the design and testing of new models as well as the incorporation of complex prior knowledge which currently is a difficult and time consuming task. Probabilistic programming languages aim to provide a formal language to specify probabilistic models in the style of computer programming and can represent any computable probability distribution as a program. In this work, we will introduce new features of Venture, a recently developed probabilistic programming language. We consider Venture the most compelling of the probabilistic programming languages because it is the first probabilistic programming language suitable for general purpose use (Mansinghka et al., 2014). Venture comes with scalable performance on hard problems and with a general purpose inference engine. The inference engine deploys Markov Chain Monte Carlo (MCMC) methods (for an introduction, see Andrieu et al. (2003)). MCMC lends itself to models with complex structures such as probabilistic programs or hierarchical Bayesian non-parametric models since they can provide a vehicle to express otherwise intractable integrals necessary for a fully Bayesian representation. MCMC is scalable, often distributable and also compositional. That is, one can arbitrarily chain MCMC kernels to infer over several hierarchically connected or nested models as they will emerge in probabilistic programming.

One very powerful model yet unseen in probabilistic programming languages are Gaussian Processes (GPs). GPs are gaining increasing attention for representing unknown functions by posterior probability distributions in various fields such as machine learning, signal processing, computer vision and bio-medical data analysis. Making GPs available in probabilistic programming is crucial to allow a language to solve a wide range of problems. Hard problems include but are not limited

054 to hierarchical prior construction (Neal, 1997), Bayesian Optimization Snoek et al. (2012) and sys-
055 tems for inductive learning of symbolic expressions such as the one introduced in the Automated
056 Statistician project Duvenaud et al. (2013); Lloyd et al. (2014). Learning such symbolic expressions
057 is a hard problem that requires careful design of approximation techniques since standard inference
058 method do not apply.

059 In the following, we will present `gpmem` as a novel probabilistic programming technique that solves
060 such hard problems. `gpmem` introduces a statistical alternative to standard memoization. Our con-
061 tribution is threefold:

- 063 • we introduce an efficient implementation of `gpmem` in form of a self-caching wrapper that
064 remembers previously computed values;
- 065 • we illustrate the statistical emulator that `gpmem` produces and how it improves with every
066 data-point that becomes available; and
- 067 • we show how one can solve hard problems of state-of-the-art machine learning related to
068 GP using `gpmem` in a Bayesian fashion and with only a few lines of Venture code.

070 We evaluate the contribution on problems posed by the GP community using real world and syn-
071 thetic data by assessing quality in terms of posterior distributions of symbolic outcome and in terms
072 of the residuals produced by our probabilistic programs. The paper is structured as follows, we will
073 first provide some background on memoization. We will explain programming in Venture and pro-
074 vide a brief introduction to GPs. We introduce `gpmem` and its use in probabilistic programming and
075 Bayesian modeling. Finally, we will show how we can apply `gpmem` on problems of causally struc-
076 tured hierarchical priors for hyper-parameter inference, structure discovery for Gaussian Processes
077 and Bayesian Optimization including experiments with real world and synthetic data.

079 2 Background

081 2.1 Memoization

083 Memoization is the idea of computing the value of a function given an input once and then store
084 it. If the function is called again at the same input, we return the stored value without evaluating
085 the function body again. Research on the Church language (Goodman et al., 2008) pointed out that
086 although memoization does not change the semantics of a program for deterministic programs it
087 does for stochastic ones. This implies that one can stochastically decide on each function evaluation
088 whether to use a stored value or a new function evaluation. Here, memoization induces random
089 world semantics (Poole, 1993; Sato, 1995) over the probabilistic program: we interpret a specific
090 mapping from function input to function output as a single possible world (McAllester et al., 2008).
091 This formulation lends it self to computation with GPs, as we will learn below.

092 2.2 Venture

094 Venture is a compositional language for custom inference strategies that comes with a Scheme- and
095 Java-Script-like front-end syntax. Its implementation is based on on three concepts. (i) stochas-
096 tic procedure interfaces that specify and encapsulate random variables, analogously to conditional
097 probability tables in a Bayesian network; (ii) probabilistic execution traces that represent execution
098 histories and capture conditional dependencies; and (iii) scaffolds that partition execution histories
099 and factor global inference problems into sub-problems. These building blocks provide a powerful
100 way to represent probability distributions; some of which cannot be expressed with density func-
101 tions. For the purpose of this work the most important Venture directives that operate on these
102 building blocks to understand are `ASSUME`, `OBSERVE`, `SAMPLE` and `INFER`. `ASSUME` induces
103 a hypothesis space for (probabilistic) models including random variables by binding the result of an
104 expression to a symbol. `SAMPLE` simulates a model expression and returns a value. `OBSERVE`
105 adds constraints to model expressions. `INFER` instructions incorporate observations and cause Ven-
ture to find a hypothesis that is probable given the data.

107 `INFER` is most commonly done by deploying the Metropolis-Hastings algorithm (MH) (Metropolis
et al., 1953). Many algorithms used in the MCMC world can be interpreted as special cases of

MH (Andrieu et al., 2003). We can outline the MH algorithm as follows. For T steps we sample x^* from a proposal distribution q :

$$x^* \sim q(x^* | x^{(t)}) \quad (1)$$

which we accept $(x^{t+1} \leftarrow x^*)$ with ratio:

$$\alpha = \min\left\{1, \frac{p(x^*)q(x^t | x^*)}{p(x^{(t)})q(x^* | x^t)}\right\} \quad (2)$$

Venture implements an MH transition operator for probabilistic execution traces.

2.3 Gaussian Processes

In the following, we will introduce GP related theory and notations. We will exclusively work on two variable regression problems. Let the data be real-valued scalars $\{x_i, y_i\}_{i=1}^n$ (complete data will be denoted by column vectors \mathbf{x}, \mathbf{y}). GPs present a non-parametric way to express prior knowledge on the space of possible functions f that we assume to have generated the data. f is assumed latent and the GP prior is given by a multivariate Gaussian $f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(x_i, x'_i))$, where $m(\mathbf{x})$ is a function of the mean of all functions that map to y_i at x_i and $k(x_i, x'_i)$ is a kernel or covariance function that summarizes the covariance of all functions that map to y_i at x_i . We can absorb the mean function into the covariance function so without loss of generality we can set the mean to zero. The marginal likelihood can be expressed as:

$$p(\mathbf{y} | \mathbf{x}) = \int p(\mathbf{y} | \mathbf{f}, \mathbf{x}) p(\mathbf{f} | \mathbf{x}) d\mathbf{f} \quad (3)$$

where the prior is Gaussian $\mathbf{f} | \mathbf{x} \sim \mathcal{N}(0, k(\mathbf{x}, \mathbf{x}'))$. We can sample a vector of unseen data from the predictive posterior with

$$\mathbf{y}^* \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \quad (4)$$

for a zero mean prior GP with a posterior mean of:

$$\boldsymbol{\mu} = \mathbf{K}(\mathbf{x}, \mathbf{x}^*) \mathbf{K}(\mathbf{x}^*, \mathbf{x}^*)^{-1} \mathbf{y} \quad (5)$$

and covariance

$$\boldsymbol{\Sigma} = \mathbf{K}(\mathbf{x}, \mathbf{x}) + \mathbf{K}(\mathbf{x}, \mathbf{x}^*) \mathbf{K}(\mathbf{x}^*, \mathbf{x}^*)^{-1} \mathbf{K}(\mathbf{x}^*, \mathbf{x}). \quad (6)$$

\mathbf{K} is a covariance function. The log-likelihood is defined as:

$$\log P(\mathbf{y} | \mathbf{X}) = -\frac{1}{2}\mathbf{y}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} - \frac{1}{2} \log |\mathbf{K} + \sigma^2 \mathbf{I}| - \frac{n}{2} \log 2\pi \quad (7)$$

with n being the number of data-points and sigma the independent observation noise. Both log-likelihood and predictive posterior can be computed efficiently in a Venture SP with an algorithm that resorts to Cholesky factorization(Rasmussen and Williams, 2006, chap. 2) resulting in a computational complexity of $\mathcal{O}(n^3)$ in the number of data-points.

The covariance function covers general high-level properties of the observed data such as linearity, periodicity and smoothness. The most widely used type of covariance function is the squared exponential covariance function:

$$k(x, x') = \sigma^2 \exp\left(-\frac{(x - x')^2}{2\ell^2}\right) \quad (8)$$

where σ and ℓ are hyper-parameters. σ is a scaling factor and ℓ is the typical length-scale. Smaller variations can be achieved by adapting these hyper-parameters.

Larger variations are achieved by changing the type of the covariance function structure. Note that covariance function structures are compositional. We can add covariance functions if we want to model globally valid structures

$$k_3(x, x') = k_1(x, x') + k_2(x, x') \quad (9)$$

and we can multiply covariance functions if the data is best explained by local structure

$$k_4(x, x') = k_1(x, x') \times k_2(x, x'); \quad (10)$$

both, k_3 and k_4 are valid covariance function structures.

162 **3 Venture GPs**
 163

164 Given a stochastic process that implements the GP algebra above we can implement
 165 a GP sampler (4) to perform GP inference in a few lines of code. We can express simple GP smoothing with fixed hyper-parameters or a prior on hyper-
 166 parameters and perform MH on it while allowing users to custom design covariance
 167 functions. Throughout the paper, we will use the Scheme-like front-end syntax.
 168

```

1 [ASSUME l (gamma 1 3)] ∈ {hyper-parameters}
2 [ASSUME sf (gamma 1 3)] ∈ {hyper-parameters}
3
4  $k(x, x') := \sigma^2 \exp\left(-\frac{(x-x')^2}{2\ell^2}\right)$ 
5
6 [ASSUME f VentureFunction(k, σ, ℓ) ]
7 [ASSUME SE make-se (apply-function f l sf) ]
8 [ASSUME (make-gp 0 SE) ]
9
10 [SAMPLE GP (array 1 2 3)] % Prior
11 [OBSERVE GP D]
12 [SAMPLE GP (array 1 2 3)]
13 [INFER (MH {hyper-parameters} one 100) ]
14 [SAMPLE GP (array 1 2 3)] % Posterior

```

181 Listing 1: Bayesian GP Smoothing
 182

183 The first two lines depict the hyper-parameters. We tag both of them to belong to the set {hyper-
 184 parameters}. Every member of this set belongs to the same inference scope. This scope controls the
 185 application of the inference procedure used. In this paper, we use MH throughout. Each scope is
 186 further subdivided into blocks that allow to do block-proposals. In the following we omit the block
 187 notation for readability, since we always choose the block of a certain scope at random.
 188

189 The ASSUME directives describe the assumptions we make for the GP model, we assume the hyper-
 190 parameters l and sf (corresponding to ℓ, σ) to be 1 and 2. The squared exponential covariance
 191 function can be defined outside the Venture code with foreign conventional programming languages,
 192 e.g. Python. In that way, the user can define custom covariance functions without being restricted to
 193 the most common ones. We then integrate the foreign function into Venture as VentureFunction. In
 194 the next line this function is associated with the hyper-parameters. Finally, we assume a Gaussian
 195 Process SP with a zero mean and the previously assumed squared exponential covariance function.
 196

197 In the case where hyper-parameters are unknown they can be found deterministically by optimizing
 198 the marginal likelihood using a gradient based optimizer. Non-deterministic, Bayesian representations
 199 of this case are also known (Neal, 1997).

200 We have already implemented this in listing 1. We draw the hyper-parameters from a Γ -prior for
 201 a Bayesian treatment of hyper-parameters. This is simple using the build in stochastic procedure
 202 that simulates drawing samples from a gamma distribution. The program gives rise to a Bayesian
 203 representation of GPs, which we will explore in the following.

204 **3.1 A Bayesian interpretation**
 205

206 **3.1.1 Data modelling as a special case of `gpmem`**

207 From the standpoint of computation, a data set of the form $\{(x_i, y_i)\}$ can be thought of as a function
 208 $y = f_{\text{restr}}(x)$, where f_{restr} is restricted to only allow evaluation at a specific set of inputs x . Modelling
 209 the data set with a GP then amounts to trying to learn a smooth function f_{emu} (“emu” stands for
 210 “emulator”) which extends f to its full domain. Indeed, if f_{restr} is a foreign procedure made available
 211 as a black-box to Venture, whose secret underlying source code is:

```

212 def f_restr(x):
213     if x in D:
214         return D[x]
215     else:
216         raise Exception('Illegal input')

```

216 Then the OBSERVE code in Listing 1 can be rewritten using gpmem as follows (where here the data
 217 set D has keys $x[1], \dots, x[n]$):

```

218 [ASSUME (list f_compute f_emu) (gpmem f_restr) ]
219 for i=1 to n:
220   [PREDICT (f_compute x[i])]
221   [INFER {MH {hyper-parameters} one 100} ]
222   [SAMPLE (f_emu (array 1 2 3))]
```

223 This rewriting has at least two benefits: (i) readability (in some cases), and (ii) amenability to active
 224 learning. As to (i), the statistical code of creating a Gaussian process is replaced with a memoization-
 225 like idiom, which will be more familiar to programmers. As to (ii), when using gpmem, it is quite
 226 easy to decide incrementally which data point to sample next: for example, the loop from $x[1]$ to
 227 $x[n]$ could be replaced by a loop in which the next index i is chosen by a supplied decision rule.
 228 In this way, we could use gpmem to perform online learning using only a subset of the available
 229 data.

233 3.1.2 The efficacy of learning hyperparameters

236 The probability of the hyper-parameters of a GP with assumptions as above and given covariance
 237 function structure \mathbf{K} can be described as:

$$241 P(\boldsymbol{\theta} | \mathbf{D}, \mathbf{K}) = \frac{P(\mathbf{D} | \boldsymbol{\theta}, \mathbf{K})P(\boldsymbol{\theta} | \mathbf{K})}{P(\mathbf{D} | \mathbf{K})}. \quad (11)$$

246 Let the \mathbf{K} be the sum of a smoothing and a white noise (WN) kernel. For this case, Neal suggested
 247 the problem of outliers in data as a use-case for a hierarchical Bayesian treatment of Gaussian
 248 processes (1997)¹. The work suggests a hierarchical system of hyper-parameterization (Fig. 1a).
 249 Here, we draw hyper-parameters from a Γ distributions:

$$254 \ell^{(t)} \sim \Gamma(\alpha_1, \beta_1), \sigma^{(t)} \sim \Gamma(\alpha_2, \beta_2) \quad (12)$$

258 and in turn sample the α and β from Γ distributions as well:

$$263 \alpha_1^{(t)} \sim \Gamma(\alpha_\alpha^1, \beta_\alpha^1), \alpha_2^{(t)} \sim \Gamma(\alpha_\alpha^2, \beta_\alpha^2), \dots \quad (13)$$

268 ¹In (Neal, 1997) the sum of an SE plus a constant kernel is used. We stick to the WN kernel for illustrative
 269 purposes.

```

270 Assuming the covariance structure is an additive comprised of a smoothing and a white noise
271 kernel, one can represent this kind of model using gpmem with only a few lines of code:
272
1 [ASSUME alpha (mem (lambda (i) (gamma 1 3 )))] ∈ {hyper-parameters-Γ}
2 [ASSUME beta (mem (lambda (i) (gamma 1 3 )))] ∈ {hyper-parameters-Γ}
273
3
274
4 [ASSUME l (gamma (alpha 1) (beta 1))] ∈ {hyper-parameters}
5 [ASSUME sf (gamma (alpha 2) (beta 2))] ∈ {hyper-parameters}
6 [ASSUME sigma (uniform 0 5 )] ∈ {hyper-parameters} % Fig. 2
275
7 % above: structured prior, Fig. 1a
276
8
277
9  $k_{SE}(x, x') := \theta^2 \exp\left(-\frac{(x-x')^2}{2\ell^2}\right)$ 
10  $k_{WN}(x, x') := \sigma^2 \delta_{x,x'}$ 
278
11
279
12 [ASSUME kSE VentureFunction(kSE, θ, ℓ) ]
280 [ASSUME kWN VentureFunction(kWN, σ) ]
281
282
283
284
285
286
287 [ASSUME (list f_compute f_emu) (gpmem f_restr (function-plus SE WN) )]
288 [SAMPLE (f_emu (array 1 2 3))] % prior, Fig. 1b
289
290
291
292
293
294
295 [INFER (REPEAT 100
296   (DO (MH {hyper-parameters} one 2)
297     (MH {hyper-parameters-Γ} one 2) ))
298 [SAMPLE (f_emu (array 1 2 3))] % posterior , Fig. 1d
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323

```

Listing 2: Hierarchical GP Smoothing

Neal provides a custom inference algorithm setting and evaluates it using the following synthetic data problem. Let f be the underlying function that generates the data:

$$f(x) = 0.3 + 0.4x + 0.5 \sin(2.7x) + \frac{1.1}{(1+x^2)} + \eta \quad \text{with } \eta \sim \mathcal{N}(0, \sigma) \quad (14)$$

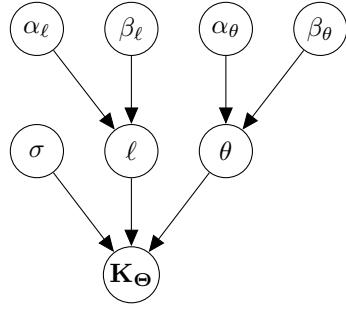
We synthetically generate outliers by setting $\sigma = 0.1$ in 95% of the cases and to $\sigma = 1$ in the remaining cases. gpmem can capture the true underlying function within only 100 MH steps on the hyper-parameters to get a good approximation for their posterior (see Fig. 1). Note that Neal devices an additional noise model and performs large number of Hybrid-Monte Carlo and Gibbs steps. We illustrate the hyper-parameter by showing the shift of the distribution on the noise parameter σ (Fig. 2). We see that gpmem learns the posterior distribution well, the posterior even exhibits a bimodal histogram when sampling σ 100 times reflecting the two modes of data generation, that is normal noise and outliers².

3.1.3 Broader applicability of gpmem

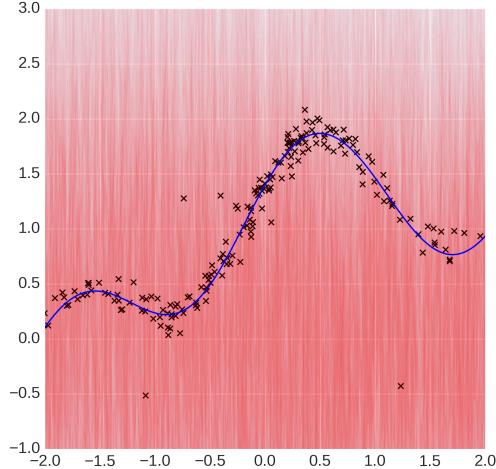
More generally, gpmem is relevant not just when a data set is available, but also whenever we have at hand a function f_{restr} which is expensive or impractical to evaluate many times. gpmem allows us to model f_{restr} with a GP-based emulator f_{emu} , and also to use f_{emu} during the learning process to choose, in an online manner, an effective set of probe points $\{x_i\}$ on which to use our few evaluations of f_{restr} . This idea is illustrated in detail in Section 4. Before doing this, we will illustrate another benefit of having a probabilistic programming apparatus for GP modelling: the linguistically unified treatment of inference over structure and inference over parameters. This unification makes interleaved joint inference over structure and parameters very natural, and allows us to give a short, elegant description of what it means to “learn the covariance function,” both in prose and in code.

²For this pedagogical example we have increased the probability for outliers in the data generation slightly from 0.05 to 0.2

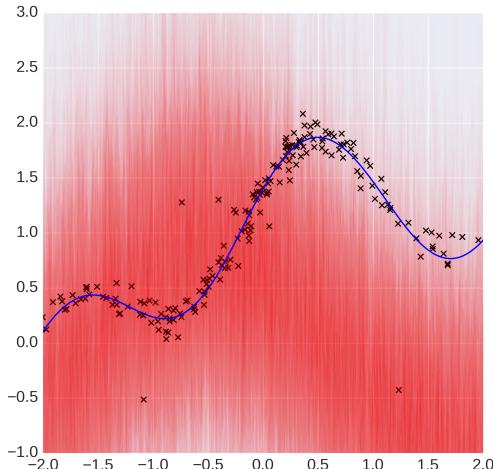
324
 325
 326
 327
 328
 329
 330
 331
 332
 333
 334
 335
 336
 337
 338
 339
 340
 341
 342
 343
 344
 345
 346



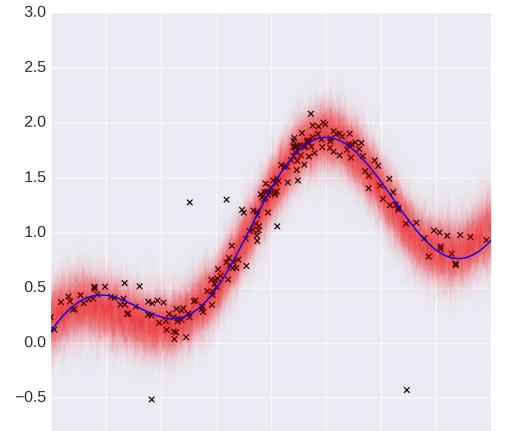
(a) Hierarchical Prior



(b) Prior Inference



(c) Observed



(d) Inferred

367
 368
 369
 370
 371
 372
 373
 374
 375
 376
 377

Figure 1: (a) depicts the hierarchical structure of the hyper-parameter as constructed in the work by Neal as a Bayesian Network. (b)-(d) shows gpmem on Neal's example. We see that prior renders functions all over the place (a). After gpmem observes a some data-points an arbitrary smooth trend with a high level of noise is sampled. After running inference on the hierarchical system of hyper-parameters we see that the posterior reflects the actual curve well. Outliers are treated as such and do not confound the GP.

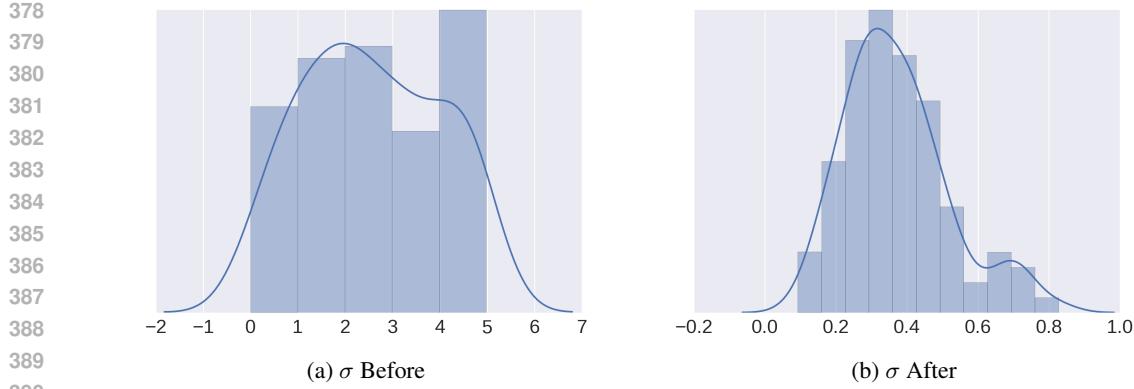


Figure 2: Hyper-parameter inference on the parameter of the noise kernel. We show 100 samples drawn from the distribution on σ . One can clearly recognise the shift from the uniform prior $\mathcal{U}(0, 5)$ to a double peak distribution around the two modes - normal and outlier.

Furthermore, the example in Section 3.2 below recovers the performance of current state-of-the-art GP-based models.

3.2 Structure Learning

The space of possible kernel composition is infinite. Combining inference over this space with the problem of finding a good parameterization that could potentially explain the observed data best poses a hard problem. The natural language interpretation of the meaning of a kernel and its composition renders this a problem of symbolic computation. Duvenaud and colleagues note that a sum of kernels can be interpreted as logical OR operations and kernel multiplication as logical AND (2013). This is due to the kernel rendering two points similar if k_1 OR k_2 outputs a high value in the case of a sum. Respectively, multiplication of two kernels results in high values only if k_1 AND k_2 have high values (see Fig. 3 exemplifies how to interpret global vs. local aspects and its symbolic analog respectively). In the following, we will refer to covariance functions that are not composite as base covariance functions.

Knowledge about the composite nature of covariance functions is not new, however, until recently, the choice and the composition of covariance functions were done ad-hoc. The Automated Statistician Project came up with an approximate search over the possible space of kernel structures (Duvenaud et al., 2013; Lloyd et al., 2014). However, a fully Bayesian treatment of this was not done before. The case where the covariance structure is not given is even more interesting. Our probabilistic programming based MCMC framework approximates the following intractable integrals of the expectation for the prediction:

$$\mathbb{E}[y^* | x^*, \mathbf{D}, \mathbf{K}] = \iint f(x^*, \boldsymbol{\theta}, \mathbf{K}) P(\boldsymbol{\theta} | \mathbf{D}, \mathbf{K}) P(\mathbf{K} | \Omega, s, n) d\boldsymbol{\theta} d\mathbf{K}. \quad (15)$$

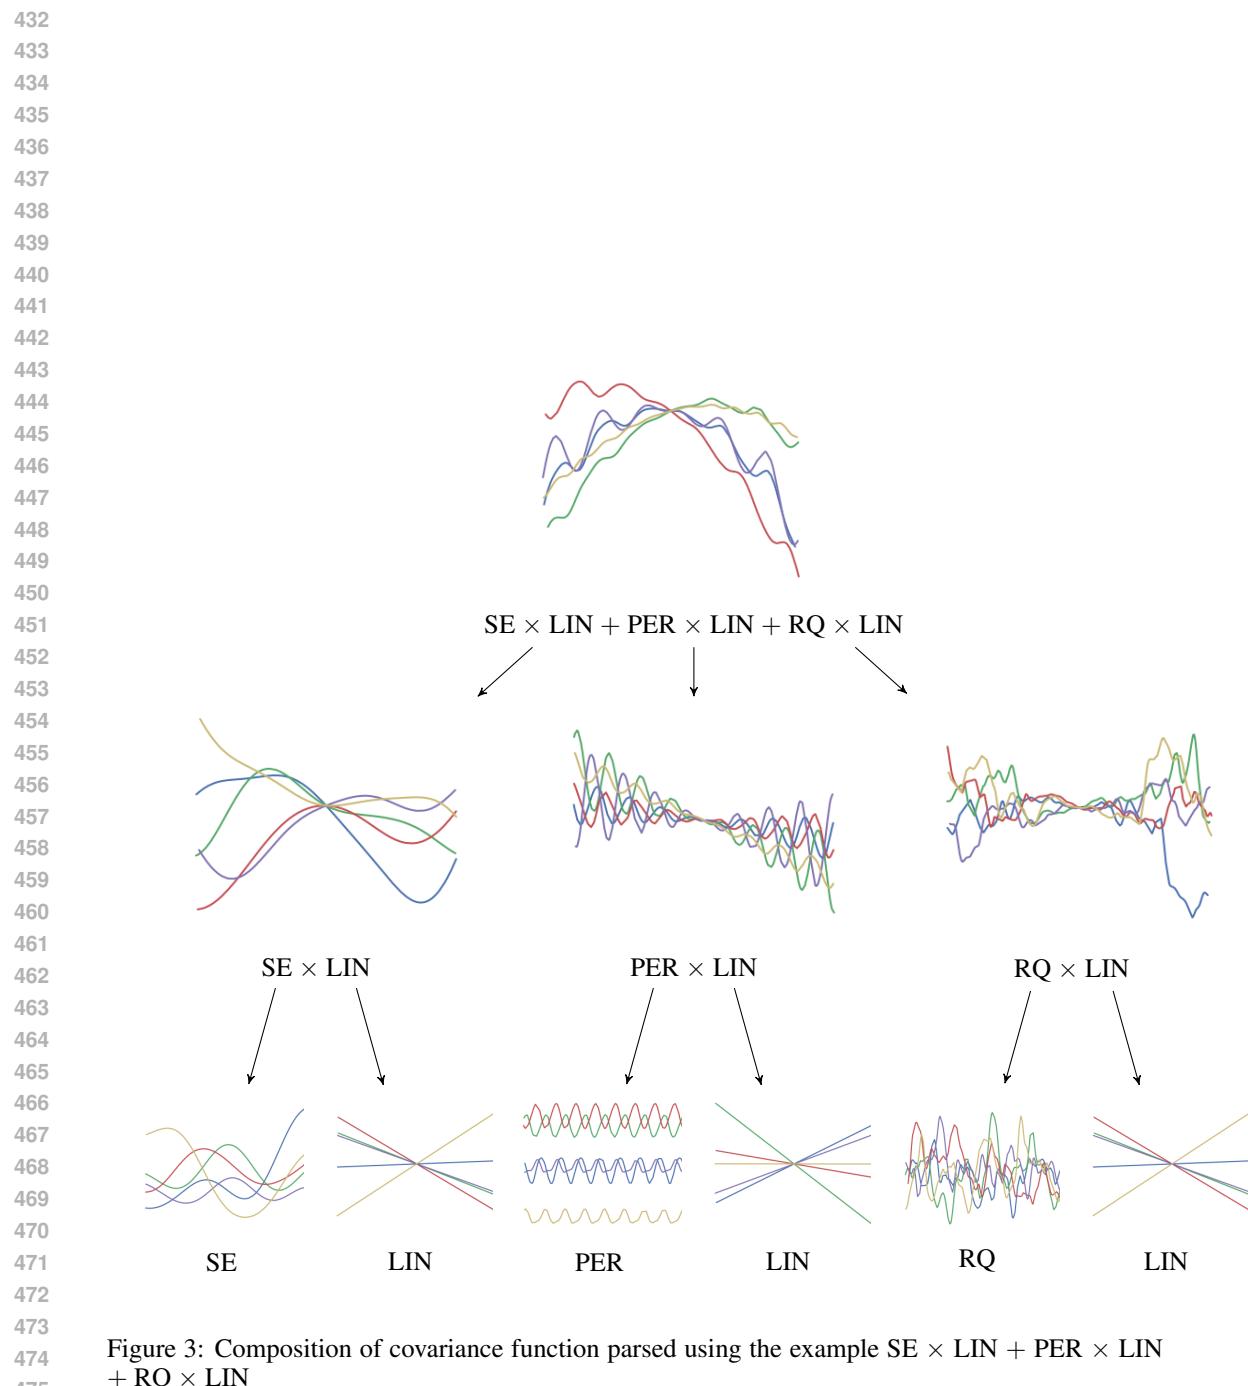
This is done by sampling from the posterior probability distribution of the hyper-parameters and the possible kernel:

$$y^* \approx \frac{1}{T} \sum_{t=1}^T f(x^* | \boldsymbol{\theta}^{(t)}, \mathbf{K}^{(t)}). \quad (16)$$

In order to provide the sampling of the kernel, we introduce a stochastic process to the SP that simulates the grammar for algebraic expressions of covariance function algebra:

$$\mathbf{K}^{(t)} \sim P(\mathbf{K} | \Omega, s, n) \quad (17)$$

Here, we start with a set of possible kernels and draw a random subset. For this subset of size n , we sample a set of possible operators that operate on the base kernels.



473 Figure 3: Composition of covariance function parsed using the example $SE \times LIN + PER \times LIN$
 474 + $RQ \times LIN$
 475

476
 477
 478
 479
 480
 481
 482
 483
 484
 485

486 The marginal probability of a kernel structure which allows us to sample is characterized by the
 487 probability of a uniformly chosen subset of the set of n possible covariance functions times the
 488 probability of sampling a global or a local structure which is given by a binomial distribution:
 489

$$P(\mathbf{K} | \boldsymbol{\Omega}, s, n) = P(\boldsymbol{\Omega} | s, n) \times P(s | n) \times P(n), \quad (18)$$

490 with
 491

$$P(\boldsymbol{\Omega} | s, n) = \binom{n}{r} p_{+ \times}^k (1 - p_{+ \times})^{n-k} \quad (19)$$

492 and
 493

$$P(s | n) = \frac{n!}{|s|!} \quad (20)$$

494 where $P(n)$ is a prior on the number of base kernels used which can sample from a discrete uniform
 495 distribution. This will strongly prefer simple covariance structures with few base kernels since
 496 individual base kernels are more likely to be sampled in this case due to (20). Alternatively, we
 497 can approximate a uniform prior over structures by weighting $P(n)$ towards higher numbers. It is
 498 possible to also assign a prior for the probability to sample global or local structures, however, we
 499 have assigned complete uncertainty to this with the probability of a flip $p = 0.5$.
 500

501 Many equivalent covariance structures can be sampled due to covariance function algebra and equiv-
 502 alent representations with different parameterization (Lloyd et al., 2014). Certain covariance func-
 503 tions can differ in terms of the hyper-parameterization but can be absorbed into a single covariance
 504 function with a different parameterization. To inspect the posterior of these equivalent structures
 505 we convert each kernel expression into a sum of products and subsequently simplify. Rules for this
 506 simplification can be found in appendix B.
 507

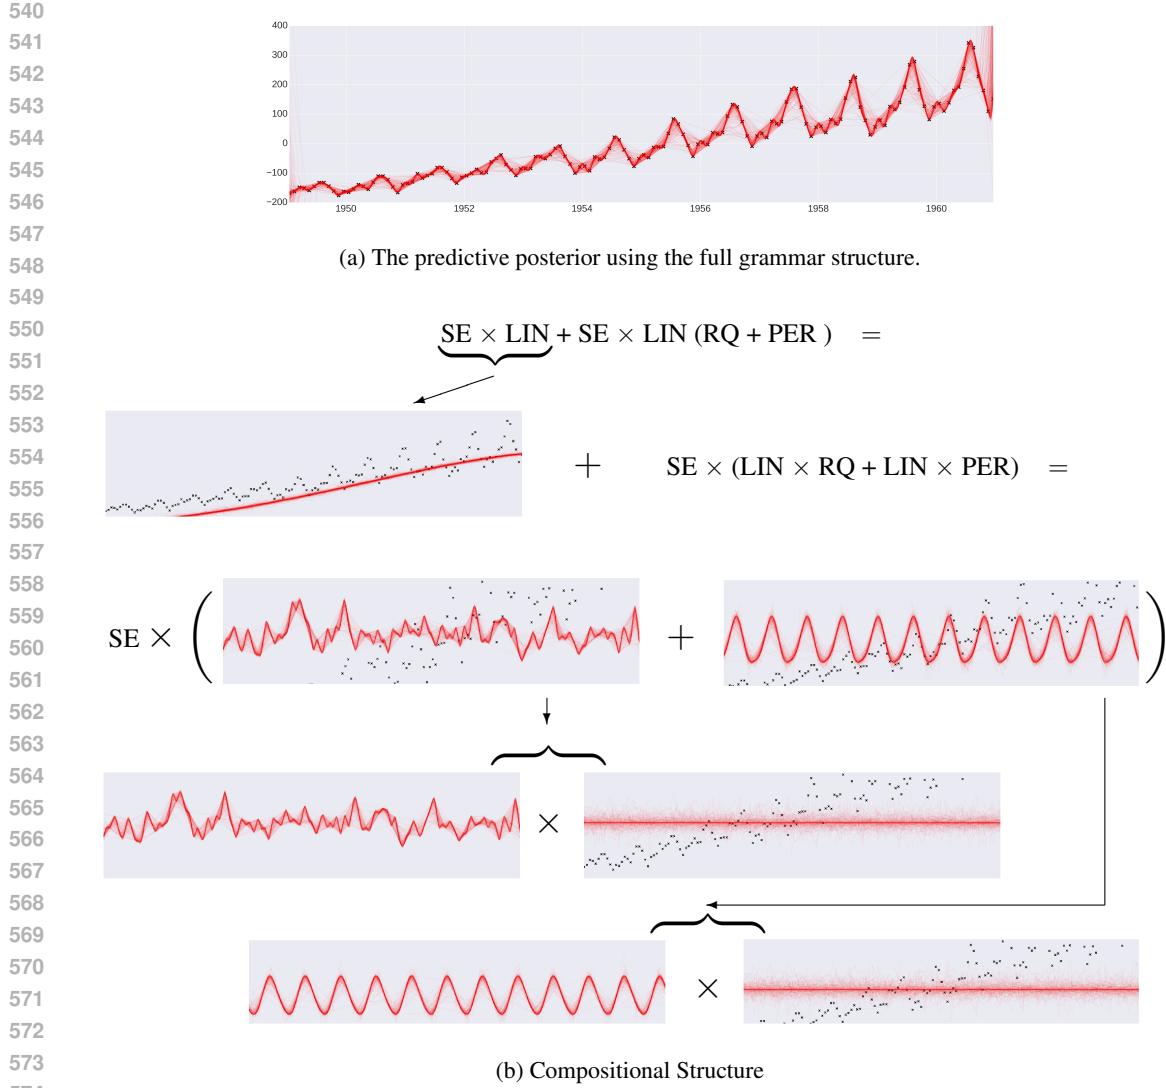
508 For reproducing results from the Automated Statistician Project in a Bayesian fashion we first define
 509 a prior on the hypothesis space. Note that, as in the implementation of the Automated Statistician,
 510 we upper-bound the complexity of the space of covariance functions we want to explore. We also
 511 put vague priors on hyper-parameters.
 512

```

1 [ASSUME base_kernels (list K1,K2,...,Kn) ] % defined as above
2 [ASSUME pn (uniform_structure n) ] % prior on the number of kernels
3 [ASSUME SK (subset base_kernels pn) ] % sampling a subset of size n
4 [ASSUME composition (lambda (l) % kernel composition
5   (if (lte ( size l) 1)
6     (first l)
7     (if (flip)
8       (func_plus (first l) (cov_compo (rest l)))
9       (func_times (first l) (cov_compo (rest l)))
10    )
11  )
12 )
13
14 [ASSUME K (composition SK) ]
15
16 [ASSUME (list f_compute f_emu) (gpmem f_restr K )]
17
18 for i=1 to n:
19   [PREDICT (f_compute x[i])] % observing with a look-up function
20
21 [INFER (REPEAT 2000 (DO
22   (MH pn one 1)
23   (MH SK one 1)
24   (MH K* one 1)
25   (MH {hyper-parameters} one 10)) ]
```

535 Listing 3: Venture Code for Bayesian GP Structure Learning
 536

537 We defined the space of covariance structures in a way allowing us to reproduce results for covari-
 538 ance function structure learning as in the Automated Statistician. This lead to coherent results, for
 539 example for the airline data set describing monthly totals of international airline passengers (Box



575 Figure 4: a) We see the predictive posterior as a result 1000 nested MH steps on the airline data
 576 set. b) depicts a decomposition of this posterior for the structures sampled by Venture.
 577 RQ is the rational quadratic covariance function. The first line shows the global trend and denotes the rest of
 578 the structure that is shown above. In the second line, the see the periodic component on the right
 579 hand side. The left hand side denotes short term deviations both multiplied by a smoothing kernel.
 580 The third and fourth lines denote how we reach the second line: both periodic and rational quadratic
 581 covariance functions are multiplied by a linear covariance function with slope zero.

582
583
584
585 et al., 1997, according to Duvenaud et al., 2013. We will elaborate the result using a sample from the
 586 posterior (Fig. 4). The sample is identical with the highest scoring result reported in previous work
 587 using a search-and-score method (Duvenaud et al., 2013) for the CO₂ data set (see Rasmussen and
 588 Williams, 2006 for a description) and the predictive capability is comparable. However, the compo-
 589 nents factor in a different way due to different parameterization of the individual base kernels.

590 We further investigated the quality of our stochastic processes by running a leave one out cross-
 591 validation to gain confidence on the posterior. This resulted in 545 independent runs of the Markov
 592 chain that produced a coherent posterior: our Bayesian interpretation of GP structure and GPs pro-
 593 duced a posterior of structures that is in line with previous results on this data set (Duvenaud et al.,
 594 2013; see Fig. 8).

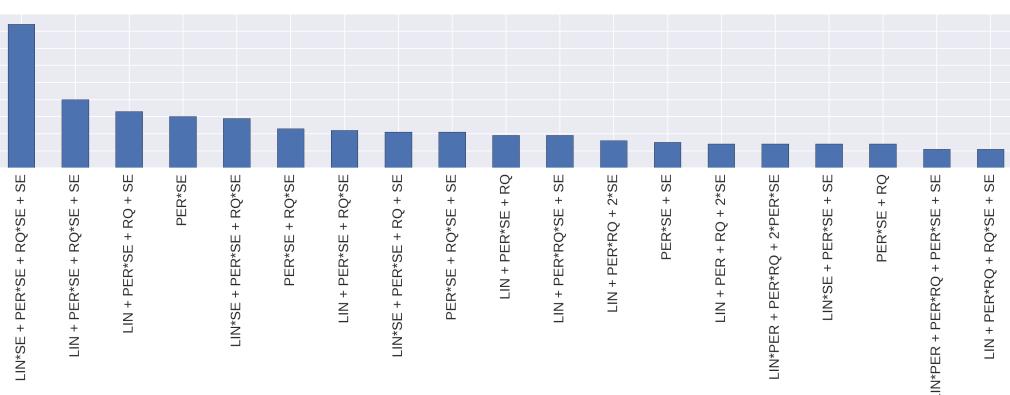


Figure 5: Posterior on structure of the CO2 data. We have cut the tail of the distribution for space reasons since the number of possible structures is large. We see the final sample of the each of the 545 chains with 2000 nested steps each. Note that Duvenaud et al. (2013) report $\text{LIN} \times \text{SE} + \text{PER} \times \text{SE} + \text{RQ} \times \text{SE}$.

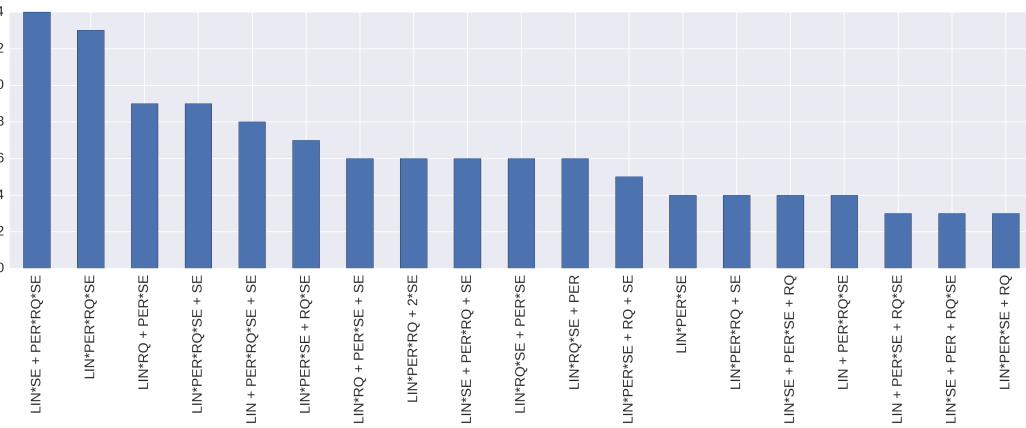


Figure 6: Posterior on structure of airline data set. We have cut the tail of the distribution for space reasons since the number of possible structures is large. We see the final sample of the each of the 144 chains with 2000 nested steps each. Note that Duvenaud et al. (2013) report $\text{LIN} \times \text{SE} + (\text{PER} + \text{RQ}) \times \text{SE} \times \text{LIN}$

We ran similar evaluation on the airline data set resulting in a similar structure to what was previously reported (Fig. 6, residuals and log-score along the Markov chain see Fig. 7).

We found the final sample of multiple runs to be most informative. This kind of Markov Chain seems to produce samples that are highly auto-correlated.

4 Bayesian Optimization

Bayesian Optimization poses the problem of finding the global maximum of an unknown function as a hierarchical decision problem (Ghahramani, 2015). Evaluating the actual function can be very expensive. For example, finding the best configuration for the learning algorithm of a large convolutional neural network implies expensive function evaluations to compare a potentially infinite number of configurations. Another common example is the example of data acquisition. For problems with large amounts of data available it may be interesting to choose certain informative data-points to evaluate a model on. In continuous domains, many Bayesian Optimization methods deploy GPs (e.g. Snoek et al., 2012).

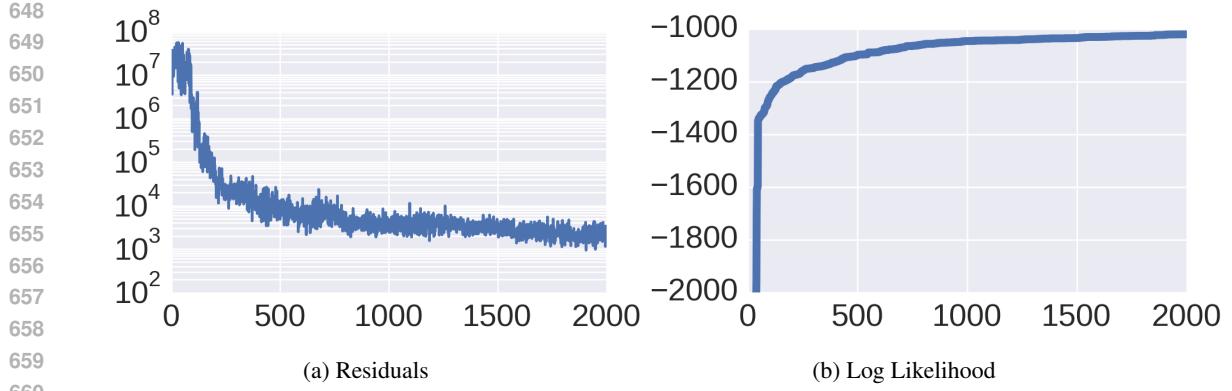


Figure 7: 2000 steps along the Markov Chain.

The hierarchical nature of Bayesian Optimization makes it an ideal application for GPs in Venture. The following Bayesian Optimization scheme is closely related to Thompson Sampling Thompson (1933), a general framework to solve exploration-exploitation problems. In our case, “exploration” is probing the true value of the function f by performing an expensive computation, and “exploitation” is using samples from a GP-based model (conditioned on the previously computed true values of f) to approximate the values of f at new points; the cost of doing so will depend on the application. In the terminology of gpmem (see Section 3.1.1), the function f_{restr} is f , and the exploration-exploitation trade-off is: given a point x , should we evaluate f_{restr} directly, thus resulting in a more accurate value and more training data for the emulator f_{emu} , or should we just sample f_{emu} for a good-enough approximation? The answer will of course depend on the application, but also note that judicious choice of when to do each will result in more useful training data for the GP model.

```

1 [ASSUME hypers ...]
2 [ASSUME se (make-squared-exponential hypers)]
3 [ASSUME (list f_compute f_emu) (gpmem f se)]
4 [ASSUME f_emu_pointwise (lambda (x) (first (f_emu (array x))))]
5 [ASSUME mc_sampler (uniform_sampler -20 20)]
6
7 for i=1 to 15:
8   [PREDICT (f_compute (mc_argmax f_emu_pointwise mc_sampler))]
9   [INFER (MH 'hypers one 50)
10
11 [INFER (collect_stats f_emu)]
```

Listing 4: Code for Bayesian optimization using gpmem. The procedure f_{compute} computes f directly, thus improving the GP model f_{emu} . ($f_{\text{emu}}_{\text{pointwise}}$ is simply a shortcut for sampling the GP model at a single point; f_{emu} is more general, allowing joint samples to be taken at any set of points.) In the loop, f_{compute} is called to compute the value of f at a new argument. The new argument, $(\text{mc_argmax } f_{\text{emu}}_{\text{pointwise}} \text{ mc_sampler})$, is a Monte Carlo estimate of the maximum pointwise sample of f_{emu} (itself a stochastic quantity), with the Monte Carlo samples being drawn in this case uniformly between -20 and 20 . After each new call to f_{compute} , the Metropolis–Hastings algorithm is used to perform inference on the hyperparameters of the covariance function in the GP model in light of the new conditioning data. Once enough calls to f_{compute} have been made (in our case we stopped at 15 calls), we can inspect the full list of probed (x, y) pairs with extract_stats . The answer to our maximization problem is simply the maximum y ; but our algorithm also learns more potentially useful information.

5 Conclusion

We have shown Venture GPs. We have introduced novel stochastic processes for a probabilistic programming language. We showed how flexible non-parametric models can be treated in Venture

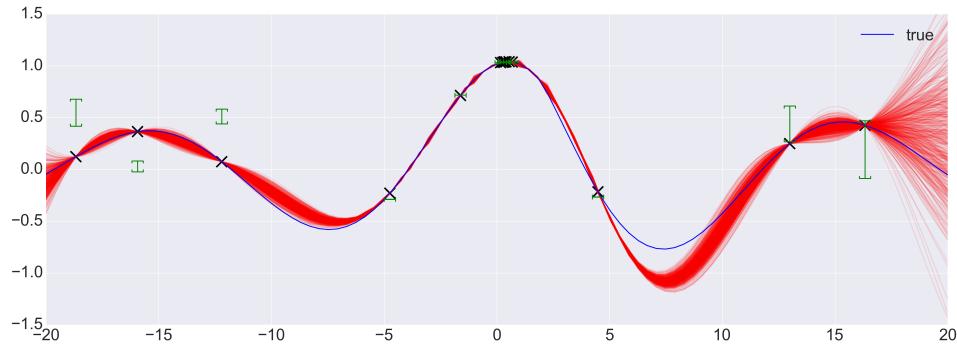


Figure 8: Bayesian Optimization. Each successive probe point x is the (stochastic) maximum of a GP-based emulator conditioned on the values of the previously probed points. In the figure, each probe point x is marked with an \times , and a vertical green bar is drawn showing the mean \pm one standard deviation of the “leave-one-out” distribution—the distribution that would arise from the same covariance function if all marked points *except* x had been probed. Note that there are many probe points near the true maximum, and the uncertainty is quite low. Also note that probed points far away from the true maximum tend to be points at which the uncertainty is high.

in only a few lines of code. We evaluated our contribution on a range of hard problems for state-of-the-art Bayesian non-parametrics. Venture GPs showed competitive performance in all of them.

Appendix

A Covariance Functions

SE and WN are defined in the text above, for completeness we will introduce the covariance:

$$k_{LIN}(x, x') = \theta(xx') \quad (21)$$

$$k_{PER}(x, x') = \theta \exp\left(\frac{2 \sin^2(\pi(x - x')/p)}{\ell^2}\right) \quad (22)$$

$$k_{RQ}(x, x') = \theta \left(1 + \frac{(x - x')^2}{2\alpha\ell^2}\right)^{-\alpha} \quad (23)$$

B Covariance Simplification

1	SE \times SE	\rightarrow SE
2	{SE, PER, C, WN} \times WN	\rightarrow WN
3	LIN + LIN	\rightarrow LIN
4	{SE, PER, C, WN, LIN} \times C	\rightarrow {SE, PER, C, WN, LIN}

Listing 5: Grammar to simplify expressions

$$\theta(xx') = \theta_1(xx') + \theta_2(xx') \quad (24)$$

756 **References**
757

- 758 Andrieu, C., De Freitas, N., Doucet, A., and Jordan, M. I. (2003). An introduction to mcmc for
759 machine learning. *Machine learning*, 50(1-2):5–43.
760 Box, G. E., Jenkins, G. M., and Reinsel, G. C. (1997). *Time series analysis: forecasting and control*.
761 Duvenaud, D., Lloyd, J. R., Grosse, R., Tenenbaum, J., and Ghahramani, Z. (2013). Structure
762 discovery in nonparametric regression through compositional kernel search. In *Proceedings of*
763 *the 30th International Conference on Machine Learning (ICML-13)*, pages 1166–1174.
764 Ghahramani, Z. (2015). Probabilistic machine learning and artificial intelligence. *Nature*,
765 521(7553):452–459.
766 Goodman, N. D .and Mansinghka, V. K., Roy, D., Bonawitz, K., and Tenenbaum, J. (2008). Church:
767 A language for generative models. In *Proceedings of the 24th Conference on Uncertainty in*
768 *Artificial Intelligence, UAI 2008*, pages 220–229.
769 Lloyd, J. R., Duvenaud, D., Grosse, R., Tenenbaum, J., and Ghahramani, Z. (2014). Automatic
770 construction and natural-language description of nonparametric regression models. In *Twenty-*
771 *Eighth AAAI Conference on Artificial Intelligence*.
772 Mansinghka, V. K., Selsam, D., and Perov, Y. (2014). Venture: a higher-order probabilistic pro-
773 gramming platform with programmable inference. *arXiv preprint arXiv:1404.0099*.
774 McAllester, D., Milch, B., and Goodman, N. D. (2008). Random-world semantics and syntactic
775 independence for expressive languages. Technical report.
776 Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., and Teller, E. (1953). Equation
777 of state calculations by fast computing machines. *The Journal of Chemical Physics*, 21(6):1087–
778 1092.
779 Neal, R. M. (1997). Monte carlo implementation of gaussian process models for bayesian regression
780 and classification. *arXiv preprint physics/9701026*.
781 Poole, D. (1993). Probabilistic horn abduction and bayesian networks. *Artificial Intelligence*,
782 64(1):81–129.
783 Rasmussen, C. E. and Williams, C. K. I. (2006). *Gaussian Processes for Machine Learning (Adap-
784 tive Computation and Machine Learning)*. The MIT Press.
785 Sato, T. (1995). A statistical learning method for logic programs with distribution semantics. In *In*
786 *Proceedings of the 12th International Conference on Logic Programming*. Citeseer.
787 Snoek, J., Larochelle, H., and Adams, R. P. (2012). Practical bayesian optimization of machine
788 learning algorithms. In *Advances in Neural Information Processing Systems*, pages 2951–2959.
789 Thompson, W. R. (1933). On the likelihood that one unknown probability exceeds another in view
790 of the evidence of two samples. *Biometrika*, pages 285–294.

793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809