

000

001

002

003

004

005

006

007

008

009

010

011

012

013

014

015

016

017

018

019

020

021

022

023

024

025

026

027

028

029

030

031

032

033

034

035

036

037

038

039

040

041

042

043

044

045

046

047

048

049

050

051

052

053

# Gaussian Processes with Probabilistic Programming

## Anonymous Author(s)

Affiliation

Address

email

## Abstract

We introduce Venture GPs, a way to formulate Gaussian Processes with probabilistic programming. Gaussian Processes are flexible non-parametric models that can be applied to a broad class of problems. We represent Gaussian Processes in Venture, a Turing complete probabilistic programming language. Venture provides a compositional language with a generalized inference engine that builds on a stochastic procedure interface. This stochastic procedure interface specifies and encapsulates primitive random variables analogously conditional probability tables in Bayesian Networks. The programming language is extended with a set of stochastic processes that allow a user to formulate Gaussian Process models and to perform numerically stable inference over them while hiding the linear algebra needed for this inside the language. We show how we can extend the non-parametric model to incorporate hierarchical causal priors on model structure and hyper-parameters with only a few lines of code. We also show state-of-the-art applications of Gaussian Processes in this framework, namely structure discovery of high-level properties of Gaussian Processes, Bayesian Optimization and hyper-parameter inference. We evaluate the performance of the programs with synthetic and real world data.

## 1 Introduction

Probabilistic programming could be revolutionary for machine intelligence due to universal inference engines and the rapid prototyping for novel models (Ghahramani, 2015). Probabilistic programming languages aim to provide a formal language to specify probabilistic models in the style of computer programming and can represent any computable probability distribution as a program. In this work, we will introduce new features of Venture, a recently developed probabilistic programming language and the first probabilistic programming language suitable for general purpose use (Mansinghka et al., 2014). Venture comes with scalable performance on hard problems and with a general purpose inference engine. The inference engine is based on Markov Chain Monte Carlo (MCMC) methods (for an introduction, see Andrieu et al. (2003)). MCMC lends itself models with complex structures such as probabilistic programs or hierarchical Bayesian non-parametric models since they can provide a vehicle to express otherwise intractable integrals necessary for a fully Bayesian representation. MCMC is scalable, often distributable and also compositional. That is, one can arbitrarily chain MCMC kernels to infer over several hierarchically connected or nested models as they will emerge in probabilistic programming.

One very powerful model yet unseen in probabilistic programming languages are Gaussian Processes (GPs). GPs are gaining increasing attention for representing unknown functions by posterior probability distributions in various fields such as machine learning, signal processing, computer vision and bio-medical data analysis. In the following, we will present Gaussian Processes as a novel feature for probabilistic programming languages. Our contribution is threefold:

- we introduce a new stochastic process for GPs in a probabilistic programming language;

- we show how one can solve hard problems of state-of-the-art machine learning with only a few lines of Venture code; and
- we introduce an additional stochastic process that samples from a probabilistic context free grammar for GP covariance structure generation.

The paper is structured as follows, we will first provide some background on probabilistic programming in Venture and GPs. We will then elaborate on our new stochastic processes. Finally, we will show how we can apply those on problems of hyper-parameter inference, structure discovery for Gaussian Processes and Bayesian Optimization including experiments with real world and synthetic data.

## 2 Background

### 2.1 Venture

Venture is a compositional language for custom inference strategies that comes with a Scheme- and Java-Script-like front-end syntax. It's implementation is based on three concepts:

1. stochastic procedure interfaces that specify and encapsulate random variables, analogously to conditional probability tables in a Bayesian network;
2. probabilistic execution traces that represent execution histories and capture conditional dependencies; and
3. scaffolds that partition execution histories and factor global inference problems into sub-problems.

These building blocks provide a powerful way to represent probability distributions; some of which cannot be expressed with density functions. For the purpose of this work the most important Venture directives that operate on these building blocks to understand are ASSUME, OBSERVE, SAMPLE and INFER. ASSUME induces a hypothesis space for (probabilistic) models including random variables by binding the result of an expression to a symbol. SAMPLE simulates a model expression and returns a value. OBSERVE adds constraints to model expressions. INFER instructions incorporate observations and cause Venture to find a hypothesis that is probable given the data.

INFER is most commonly done by deploying the Metropolis-Hastings algorithm (MH) (Metropolis et al., 1953). Many algorithms used in the MCMC world can be interpreted as special cases of MH (Andrieu et al., 2003). We can outline the MH algorithm as follows. For  $T$  steps we sample  $x^*$  from a proposal distribution  $q$ :

$$x^* \sim q(x^* | x^{(t)}) \quad (1)$$

which we accept ( $x^{t+1} \leftarrow x^*$ ) with ratio:

$$\alpha = \min \left\{ 1, \frac{p(x^*)q(x^t | x^*)}{p(x^{(t)})q(x^* | x^t)} \right\} \quad (2)$$

Venture implements an MH transition operator for probabilistic execution traces.

### 2.2 Gaussian Processes

In the following, we will introduce GP related theory and notations. We will exclusively work on two variable regression problems. Let the data be real-valued scalars  $\{x_i, y_i\}_{i=1}^n$  (complete data will be denoted by column vectors  $\mathbf{x}, \mathbf{y}$ ). GPs present a non-parametric way to express prior knowledge on the space of possible functions  $f$  that we assume to have generated the data.  $f$  is assumed latent and the GP prior is given by a multivariate Gaussian  $f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(x_i, x'_i))$ , where  $m(\mathbf{x})$  is a function of the mean of all functions that map to  $y_i$  at  $x_i$  and  $k(x_i, x'_i)$  is a kernel or covariance function that summarizes the covariance of all functions that map to  $y_i$  at  $x_i$ . We can absorb the mean function into the covariance function so without loss of generality we can set the mean to zero. The marginal likelihood can be expressed as:

$$p(\mathbf{y}|\mathbf{x}) = \int p(\mathbf{y}|\mathbf{f}, \mathbf{x}) p(\mathbf{f}|\mathbf{x}) d\mathbf{f} \quad (3)$$

108 where the prior is Gaussian  $\mathbf{f}|\mathbf{x} \sim \mathcal{N}(0, k(\mathbf{x}, \mathbf{x}'))$ . We can sample a vector of unseen data from the  
 109 predictive posterior with

$$110 \quad \mathbf{y}^* \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \quad (4)$$

111 for a zero mean GP with mean:

$$113 \quad \boldsymbol{\mu} = \mathbf{K}(\mathbf{x}, \mathbf{x}^*) \mathbf{K}(\mathbf{x}^*, \mathbf{x}^*)^{-1} \mathbf{y} \quad (5)$$

114 and covariance

$$115 \quad \boldsymbol{\Sigma} = \mathbf{K}(\mathbf{x}, \mathbf{x}) + \mathbf{K}(\mathbf{x}, \mathbf{x}^*) \mathbf{K}(\mathbf{x}^*, \mathbf{x}^*)^{-1} \mathbf{K}(\mathbf{x}^*, \mathbf{x}). \quad (6)$$

116  $\mathbf{K}$  is a covariance function. The log-likelihood is defined as:

$$118 \quad \log P(\mathbf{y} | \mathbf{X}) = -\frac{1}{2} \mathbf{y}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} - \frac{1}{2} \log |\mathbf{K} + \sigma^2 \mathbf{I}| - \frac{n}{2} \log 2\pi \quad (7)$$

120 with  $n$  being the number of data-points and sigma the independent observation noise. Both log-  
 121 likelihood and predictive posterior can be computed efficiently in a Venture SP with an algorithm  
 122 that resorts to Cholesky factorization(Rasmussen and Williams, 2006, chap. 2) resulting in a com-  
 123 putational complexity of  $\mathcal{O}(n^3)$  in the number of data-points.

124 The covariance function covers general high-level properties of the observed data such as linear-  
 125 ity, periodicity and smoothness. The most widely used type of covariance function is the squared  
 126 exponential covariance function:

$$128 \quad k(x, x') = \sigma^2 \exp\left(-\frac{(x - x')^2}{2\ell^2}\right) \quad (8)$$

130 where  $\sigma$  and  $\ell$  are hyper-parameters.  $\sigma$  is a scaling factor and  $\ell$  is the typical length-scale. Smaller  
 131 variations can be achieved by adapting these hyper-parameters.

### 133 3 Venture GPs

135 Below, we see how we can express simple GP smoothing with a few lines  
 136 of Venture code while allowing users to custom design covariance func-  
 137 tions. Throughout the paper, we will use the Scheme-like front-end syntax.

139 Listing 1: GP Smoothing

```
140 [ASSUME l 1] ∈ {hyper-parameters}
141 [ASSUME sf 2] ∈ {hyper-parameters}
142
143 k(x, x') := σ² exp(-((x - x')²)/(2ℓ²))
144
145 [ASSUME f VentureFunction(k, σ, ℓ) ]
146 [ASSUME SE make-se (apply-function f l sf) ]
147 [ASSUME (make-gp 0 SE) ]
148
149 [SAMPLE GP (array 1 2 3)] % Prior
150
151 [OBSERVE GP D]
152
153 [SAMPLE GP (array 1 2 3)]
154
155 [INFER (MH {hyper-parameters} one 100) ]
156
157 [SAMPLE GP (array 1 2 3)] % Posterior
```

158 The first two lines depict the hyper-parameters. We tag both of them to belong to the set {hyper-  
 159 parameters}. Every member of this set belongs to the same inference scope. This scope controls the  
 160 application of the inference procedure used. In this paper, we use MH throughout. Each scope is  
 161 further subdivided into blocks that allow to do block-proposals. In the following we omit the block  
 notation for readability, since we always choose the block of a certain scope at random.

162 The ASSUME directives describe the assumptions we make for the GP model, we assume the hyper-  
 163 parameters 1 and sf (corresponding to  $\ell, \sigma$ ) to be 1 and 2. The squared exponential covariance  
 164 function can be defined outside the Venture code with foreign conventional programming languages,  
 165 e.g. Python. In that way, the user can define custom covariance functions without being restricted to  
 166 the most common ones. We then integrate the foreign function into Venture as VentureFunction. In  
 167 the next line this function is associated with the hyper-parameters. Finally, we assume a Gaussian  
 168 Process SP with a zero mean and the previously assumed squared exponential covariance function.

169 In the case where hyper-parameters are unknown they can be found deterministically by optimizing  
 170 the marginal likelihood using a gradient based optimizer. Non-deterministic, Bayesian representa-  
 171 tions of this case are also known (Neal, 1997). Extending the program described in listing 1 for a  
 172 Bayesian treatment of hyper-parameters is simple using the build in stochastic procedure that simu-  
 173 lates drawing samples from a gamma distribution:

174  
 175 Listing 2: Bayesian GP Smoothing

```
176 [ASSUME 1 (gamma 1 3)]
177 [ASSUME sf (gamma 1 2)]
178
179  $k(x, x') := \sigma^2 \exp\left(-\frac{(x-x')^2}{2\ell^2}\right)$ 
180
181 [ASSUME f VentureFunction ( $k, \sigma, \ell$ ) ]
182 [ASSUME SE make-se (apply-function f 1 sf) ]
183 [ASSUME (make-gp 0 SE ) ]
```

184 Larger variations are achieved by changing the type of the covariance function structure. A different  
 185 type could be a linear covariance function:  
 186

$$k(x, x') = \sigma^2(x - \ell)(x' - \ell). \quad (9)$$

187 Note that covariance function structures are compositional. We can add covariance functions if we  
 188 want to model globally valid structures  
 189

$$k_3(x, x') = k_1(x, x') + k_2(x, x') \quad (10)$$

190 and we can multiply covariance functions if the data is best explained by local structure  
 191

$$k_4(x, x') = k_1(x, x') \times k_2(x, x'); \quad (11)$$

192 both,  $k_3$  and  $k_4$  are valid covariance function structures. This leads to an infinite space of possible  
 193 structures that could potentially explain the observed data best (e.g. Fig. 1). In the following, we  
 194 will refer to covariance functions that are not composite as base covariance functions. Note that this  
 195 form of composition can be easily expressed in Venture, for example if one wishes to add a linear  
 196 and a periodic kernel:  
 197

200  
 201 Listing 3: LIN  $\times$  PER

```
202 [ASSUME 1 (gamma 1 3)]
203 [ASSUME sf (gamma 1 2)]
204 [ASSUME a (gamma 2 2)]
205
206  $k_{LIN}(x, x') = \sigma_1^2(x - \ell)(x' - \ell)$ 
207  $k_{PER}(x, x') := \sigma_2^2 \exp\left(-\frac{2 \sin^2(\pi(x-x')/\ell)}{\ell^2}\right)$ 
208
209 [ASSUME fLIN VentureFunction ( $k_{LIN}, \sigma_1$ ) ]
210 [ASSUME fPER VentureFunction ( $k_{PER}, \sigma_2, \ell, p$ ) ]
211 [ASSUME LIN (make-LIN (apply-function fLIN a)) ]
212 [ASSUME PER (make-PER (apply-function fPER 1 sf)) ]
213 [ASSUME (make-gp 0 (function-times LIN PER)) ]
```

214 Knowledge about the composite nature of covariance functions is not new, however, until recently,  
 215 the choice and the composition of covariance functions were done ad-hoc. The Automated Statisti-

cian Project came up with an approximate search over the possible space of kernel structures (Duvenaud et al., 2013; Lloyd et al., 2014). However, a fully Bayesian treatment of this was not done before.

### 3.1 A Bayesian interpretation

In the following, we will explore a Bayesian representation of GP. The probability of the hyper-parameters of a GP with assumptions as above and given covariance function structure  $\mathbf{K}$  can be described as:

$$P(\boldsymbol{\theta} | \mathbf{D}, \mathbf{K}) = \frac{P(\mathbf{D} | \boldsymbol{\theta}, \mathbf{K})P(\boldsymbol{\theta} | \mathbf{K})}{P(\mathbf{D} | \mathbf{K})}. \quad (12)$$

Neal suggested the treatment of outliers as a use-case for a Bayesian treatment of Gaussian processes (1997). He evaluates his MCMC setting using the following synthetic data problem. Let  $f$  be the underlying function that generates the data:

$$f(x) = 0.3 + 0.4x + 0.5 \sin(2.7x) + \frac{1.1}{(1+x^2)} + \eta \quad \text{with } \eta \sim \mathcal{N}(0, \sigma) \quad (13)$$

We synthetically generate outliers by setting  $\sigma = 0.1$  in 95% of the case and to  $\sigma = 1$  in the remaining cases. Venture GPs can capture the true underlying function within only 100 MH steps (see Fig. 2). Note that Neal devices an additional noise model and performs large numbe of Hybrid-Monte Carlo and Gibbs steps.

### 3.2 Structure Learning

The case where the covariance structure is not given is even more interesting. Our probabilistic programming based MCMC framework approximates the following intractable integrals of the expectation for the prediction:

$$\mathbb{E}[y^* | x^*, D, \mathbf{K}_\Omega^s] = \iint f(x^*, \boldsymbol{\theta}, \mathbf{K}) P(\boldsymbol{\theta} | \mathbf{D}, \mathbf{K}) P(\mathbf{K} | \boldsymbol{\Omega}, s, n) d\boldsymbol{\theta} d\mathbf{K}. \quad (14)$$

This is done by sampling from the posterior probability distribution of the hyper-parameters and the possible kernel:

$$y^* \approx \frac{1}{T} \sum_{t=1}^T f(x^* | \boldsymbol{\theta}^{(t)}, \mathbf{K}^{(t)}). \quad (15)$$

In order to provide the sampling of the kernel, we introduce a stochastic process to the SP that simulates the grammar for algebraic expressions of kernel algebra. Here, we start with a set of possible kernels and draw a random subset. For this subset of size  $n$ , we sample a set of possible operators that operate on the base kernels.

The marginal probability of a kernel structure which allows us to sample is characterized by the probability of a uniformly chosen subset of the set of  $n$  possible covariance functions times the probability of sampling a global or a local structure which is given by a binomial distribution:

$$P(\mathbf{K} | \boldsymbol{\Omega}, s, n) = P(\boldsymbol{\Omega} | s, n) \times P(s | n) \times P(n), \quad (16)$$

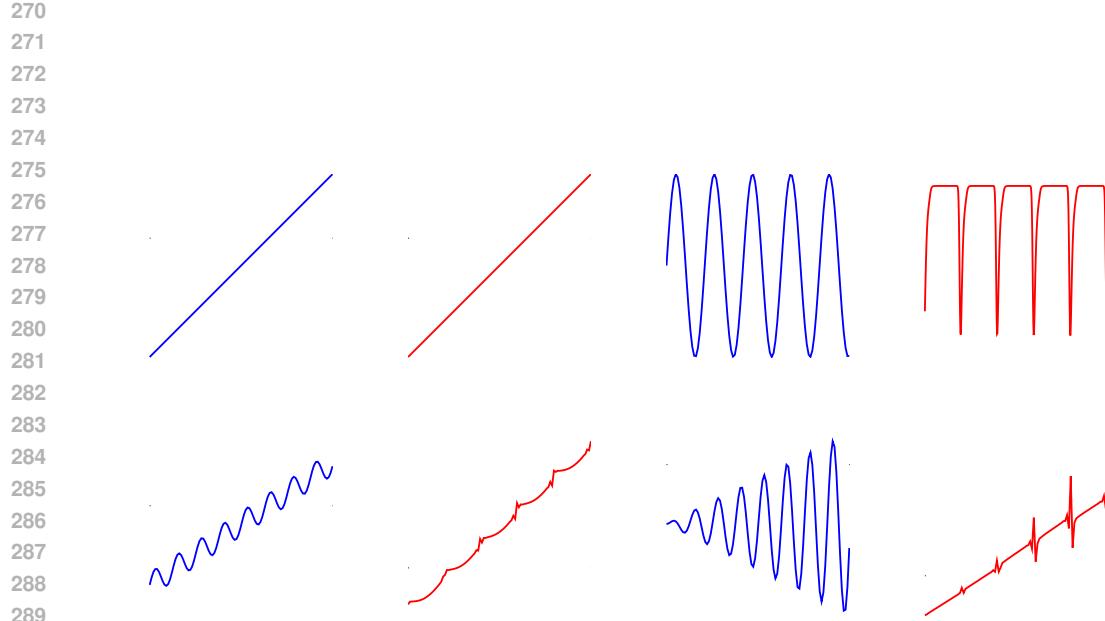
with

$$P(\boldsymbol{\Omega} | s, n) = \binom{n}{r} p_{+ \times}^k (1 - p_{+ \times})^{n-k} \quad (17)$$

and

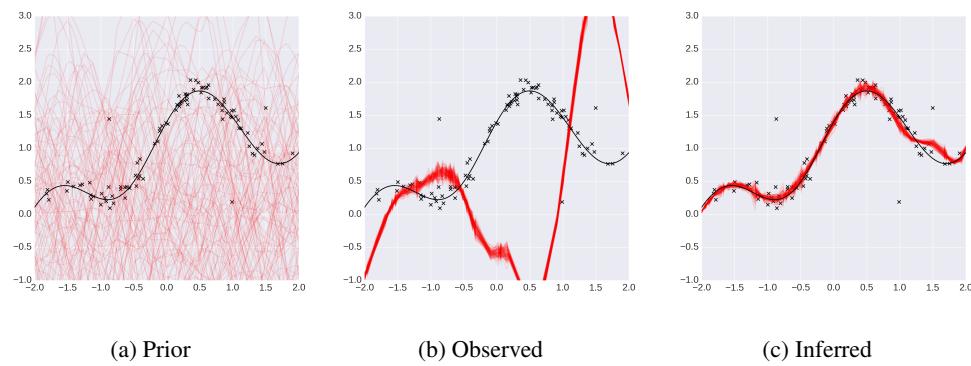
$$P(s | n) = \frac{n!}{|s|!} \quad (18)$$

where  $P(n)$  is a prior on the number of base kernels used. It is possible to also assign a prior for the probability to sample global or local priors, however, we have assigned complete uncertainty to this with the binomial  $p = 0.5$ .



290  
291  
292  
293  
294  
295  
296  
297  
298  
299  
300  
301  
302  
303  
304  
305  
306  
307  
308  
309  
310  
311  
312  
313  
314  
315

Figure 1: Composition of covariance functions (blue, left) and samples from the distribution of curves they can produce (red, right).



316  
317  
318  
319  
320  
321  
322  
323

Figure 2: Running a Venture GP on Neal's example for MCMC showing the prior, after having observed the data and after performing inference on the hyper-parameters. Note how the GP is choosing outliers to smooth instead of essential data before inference takes place.

324 Many equivalent covariance structures can be sampled due to covariance function algebra  
 325 and equivalent representations with different parameterization (Lloyd et al., 2014). Certain  
 326 covariance functions can differ in terms of the hyper-parameterization but can be  
 327 absorbed into a single covariance function with a different parameterization. To inspect  
 328 the posterior of these equivalent structures we convert each kernel expression into  
 329 a sum of products and subsequently simplify expressions using the following grammar:  
 330

331 Listing 4: Grammar to simplify expressions

332 SE × SE	→ SE
333 {SE, PER, C, WN} × WN	→ WN
334 LIN + LIN	→ LIN
335 {SE, PER, C, WN, LIN} × C	→ {SE, PER, C, WN, LIN}

336 For reproducing results from the Automated Statistician Project in a Bayesian fashion we first define  
 337 a prior on the hypothesis space. Note that, as in the implementation of the Automated Statistician,  
 338 we upper-bound the complexity of the space of covariance functions we want to explore. We also  
 339 put vague priors on hyper-parameters.  
 340

341 Listing 5: Venture Code for Bayesian GP Structure Learning

```

342 [ASSUME S (array K1,K2,...,Kn)] // (defined as above)
343 [ASSUME pn (uniform_structure n)]
344 [ASSUME S (array K1,K2,...,Kn)]
345 [ASSUME K* (grammar S pn)]
346 [ASSUME GP (make-gp 0 K*)]
347
348 [OBSERVE GP D]
349
350 [INFER (REPEAT 2000 (DO
351   (MH 10 pn one 1)
352   (MH 10 K* one 1)
353   (MH 10 {hyper-parameters} one 10)) ]
354

```

355 We defined the space of covariance structures in a way allowing us to reproduce results for covariance  
 356 function structure learning as in the Automated Statistician. This lead to coherent results, for example  
 357 for the airline data set. We will elaborate the result using a sample from the posterior (Fig. 3). The sample is identical with the highest scoring result reported in previous work using a search-and-score method (Duvenaud et al., 2013) and the predictive capability is comparable. However, the components factor in a different way due to different parameterization of the individual base kernels.

361 We further investigated the quality of our stochastic processes by running a leave one out cross-validation to gain confidence on the posterior. This resulted in 545 independent runs of the Markov chain that produced a coherent posterior: our Bayesian interpretation of GP structure and GPs produced a posterior of structures that is in line with previous results on this data set ( Duvenaud et al., 2013; see Fig. 4).

366 We found the final sample of multiple runs to be most informative. This kind of Markov Chain  
 367 seems to produce samples that are highly auto-correlated.  
 368

## 370 4 Bayesian Optimization

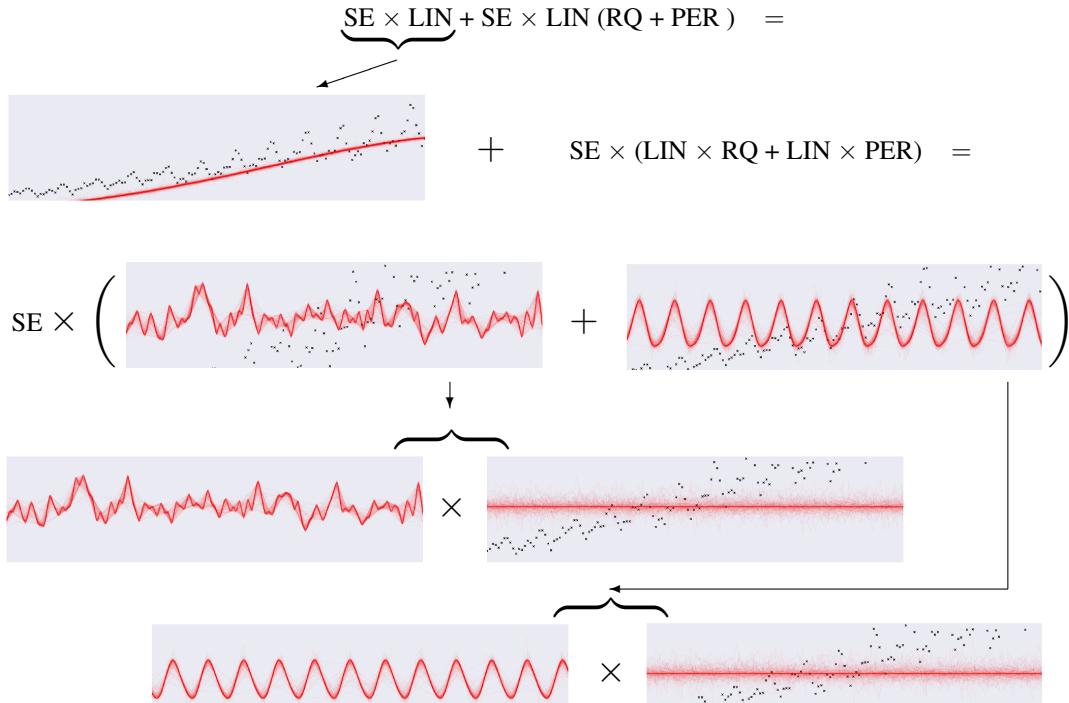
373 Bayesian Optimization poses the problem of finding the global maximum of an unknown function as a hierarchical decision problem (Ghahramani, 2015). Evaluating the actual function can be  
 374 very expensive. For example, finding the best configuration for the learning algorithm of a large  
 375 convolutional neural network implies expensive function evaluations to compare a potentially infinite  
 376 number of configurations. Another common example is the example of data acquisition. For problems with large amounts of data available it may be interested to chose certain informative data-

378  
379  
380  
381  
382  
383  
384  
385  
386  
387  
388  
389  
390



(a) The predictive posterior on given the full grammar structure.

391  
392  
393  
394  
395  
396  
397  
398  
399  
400  
401  
402  
403  
404  
405  
406  
407  
408  
409  
410  
411  
412  
413  
414  
415  
416  
417



(b) Compositional Structure

418  
419  
420  
421  
422  
423  
424  
425  
426  
427  
428  
429  
430  
431

Figure 3: a) We see the predictive posterior as an result 1000 nested MH steps on the airline data set. b) depicts a decomposition of this posterior for the structures sampled by Venture. RQ is the rational quadratic covariance function. Note that although the overall result is in line with what Duvenaud et al. (2013) report a slightly different composition is implied by the sampled parameters for the structure. The first line shows the global trend and denotes the rest of the structure that is shown above. In the second line, the see the periodic component on the right hand side. The left hand side denotes short term deviations both multiplied by a smoothing kernel. The third and fourth lines denote how we reach the second line: both periodic and rational quadratic covariance functions are multiplied by a line with slope zero.

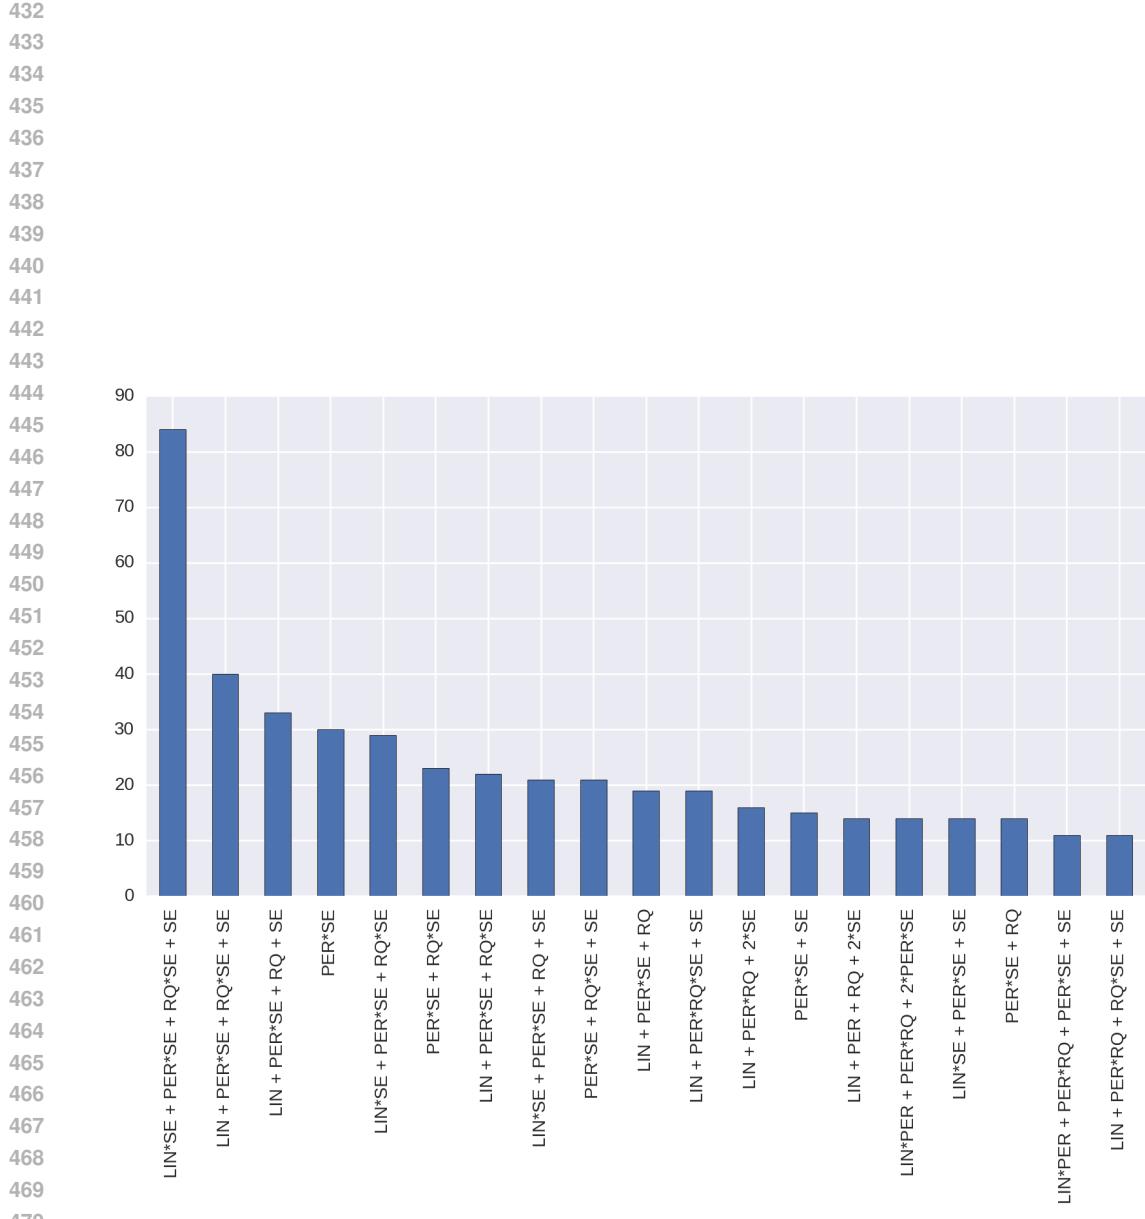


Figure 4: Posterior on structure of the CO<sub>2</sub> data. We have cut the tail of the distribution for space reasons since the number of possible structures is large. We see the final sample of the each of the 545 chains with 2000 nested steps each. Note that Duvenaud et al. (2013) report LIN × SE + PER × SE + RQ × SE.

486 points to evaluate a model on. In continuous domains, many Bayesian Optimization methods deploy  
487 GPs (e.g. Snoek et al., 2012).

488 The hierarchical nature of Bayesian Optimization makes it an ideal application for GPs in Venture.  
489 The following Bayesian Optimization scheme is closely related to Thompson Sampling Thompson  
490 (1933). Thompson Sampling is a general framework to solve exploration-exploitation problems that  
491 applies to our notion of Bayesian Optimization. We we sample a probe from the posterior  
492

$$\hat{\theta} \sim P(\hat{\theta} | \vec{x}) \quad (19)$$

493 and then optimize the expected reward by choosing an action  $a$   
494

$$a_i = \arg \max_a \mathbb{E}_{P(\dots|\theta)}[r(world_\theta(a))] \quad (20)$$

$$x_i \sim world_\theta(a_i) \quad (21)$$

## 500 References

### 501

- 502 Andrieu, C., De Freitas, N., Doucet, A., and Jordan, M. I. (2003). An introduction to mcmc for  
503 machine learning. *Machine learning*, 50(1-2):5–43.
- 504 Duvenaud, D., Lloyd, J. R., Grosse, R., Tenenbaum, J., and Ghahramani, Z. (2013). Structure  
505 discovery in nonparametric regression through compositional kernel search. In *Proceedings of  
506 the 30th International Conference on Machine Learning (ICML-13)*, pages 1166–1174.
- 507 Ghahramani, Z. (2015). Probabilistic machine learning and artificial intelligence. *Nature*,  
508 521(7553):452–459.
- 509 Lloyd, J. R., Duvenaud, D., Grosse, R., Tenenbaum, J., and Ghahramani, Z. (2014). Automatic  
510 construction and natural-language description of nonparametric regression models. In *Twenty-  
511 Eighth AAAI Conference on Artificial Intelligence*.
- 512 Mansinghka, V. K., Selsam, D., and Perov, Y. (2014). Venture: a higher-order probabilistic pro-  
513 gramming platform with programmable inference. *arXiv preprint arXiv:1404.0099*.
- 514 Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., and Teller, E. (1953). Equation  
515 of state calculations by fast computing machines. *The Journal of Chemical Physics*, 21(6):1087–  
516 1092.
- 517 Neal, R. M. (1997). Monte carlo implementation of gaussian process models for bayesian regression  
518 and classification. *arXiv preprint physics/9701026*.
- 519 Rasmussen, C. E. and Williams, C. K. I. (2006). *Gaussian Processes for Machine Learning (Adap-  
520 tive Computation and Machine Learning)*. The MIT Press.
- 521 Snoek, J., Larochelle, H., and Adams, R. P. (2012). Practical bayesian optimization of machine  
522 learning algorithms. In *Advances in Neural Information Processing Systems*, pages 2951–2959.
- 523 Thompson, W. R. (1933). On the likelihood that one unknown probability exceeds another in view  
524 of the evidence of two samples. *Biometrika*, pages 285–294.

527  
528  
529  
530  
531  
532  
533  
534  
535  
536  
537  
538  
539