

---

# Probabilistic Programming with Gaussian Process Memoization

---

Anonymous Author(s)

Affiliation  
Address  
email

## Abstract

This paper describes the *Gaussian process memoizer*, a probabilistic programming technique that uses Gaussian processes to provide a statistical alternative to memorization. Memoizing a target procedure results in a self-caching wrapper that remembers previously computed values. Gaussian process memoization additionally produces a statistical emulator based on a Gaussian process whose predictions automatically improve whenever a new value of the target procedure becomes available. This paper also introduces an efficient implementation, named `gpmem`, that can use kernels given by a broad class of probabilistic programs. The flexibility of `gpmem` is illustrated via three applications: (i) GP regression with hierarchical hyper-parameter learning, (ii) Bayesian structure learning via compositional kernels generated by a probabilistic grammar, and (iii) a bandit formulation of Bayesian optimization with automatic inference and action selection. All applications share a single 50-line Python library and require fewer than 20 lines of probabilistic code each.

## 1 Introduction

Gaussian processes are widely used tools in statistics (Barry, 1986), machine learning (Neal, 1995; Williams and Barber, 1998; Kuss and Rasmussen, 2005; Rasmussen and Williams, 2006; Damianou and Lawrence, 2013), robotics (Ferris et al., 2006), computer vision (Kemmler et al., 2013), and scientific computation (Kennedy and O’Hagan, 2001; Schneider et al., 2008; Kwan et al., 2013). They are also central to probabilistic numerics, an emerging effort to develop more computationally efficient numerical procedures, and to Bayesian optimization, a family of meta-optimization techniques that are widely used to tune parameters for deep learning algorithms (Snoek et al., 2012; Gelbart et al., 2014). They have even seen use in artificial intelligence; for example, they provide the key technology behind a project that produces qualitative natural language descriptions of time series (Duvenaud et al., 2013; Lloyd et al., 2014).

This paper describes Gaussian process memoization, a technique for integrating Gaussian processes into a probabilistic programming language, and demonstrates its utility by re-implementing and extending state-of-the-art applications of the GP. Memoization, typically implemented by a procedure called `mem()`, is a classic higher-order programming technique in which a procedure is augmented with an input-output cache that is checked each time before the function is invoked. This prevents unnecessary recomputation, potentially saving time at the cost of increased storage requirements. Gaussian process memoization, implemented by the `gpmem()` procedure, generalizes this idea to include a statistical emulator that uses previously computed values as data in a statistical model that can cheaply forecast probable outputs. The covariance function for the Gaussian process is also allowed to be an arbitrary probabilistic program.

This paper presents three applications of `gpmem`: (i) a replication of (Neal, 1997) results on outlier rejection via hyper-parameter inference; (ii) a fully Bayesian extension to the Automated Statis-

054  
055  
056  
057  
058  
059

tician project; and (iii) an implementation of Bayesian optimization via Thompson sampling. The first application can in principle be replicated in several other probabilistic languages embedding the proposal that is described in this paper. The remaining two applications rely on distinctive capabilities of Venture: support for fully Bayesian structure learning and language constructs for inference programming. All applications share a single 50-line Python library and require fewer than 20 lines of probabilistic code each.

060

## 061    2 Gaussian Processes Memoization

062

063  
064  
065  
066  
067  
068  
069

Let the data be pairs of real-valued scalars  $\{(x_i, y_i)\}_{i=1}^n$  (complete data will be denoted by column vectors  $\mathbf{x}, \mathbf{y}$ ). Gaussian Process (GP)s present a non-parametric way to express prior knowledge on the space of possible functions  $f$  modeling a regression relationship. Formally, a GP is an infinite-dimensional extension of the multivariate Gaussian distribution. Often one assumes the values  $\mathbf{y}$  are noisily measured, that is, one only sees the values of  $\mathbf{y}_{\text{noisy}} = \mathbf{y} + \mathbf{w}$  where  $\mathbf{w}$  is Gaussian white noise with variance  $\sigma_{\text{noise}}^2$ . In that case, the log-likelihood of a GP is

$$070 \quad \log p(\mathbf{y}_{\text{noisy}} | \mathbf{x}) = -\frac{1}{2}\mathbf{y}^\top(\Sigma + \sigma_{\text{noise}}^2 \mathbf{I})^{-1}\mathbf{y} - \frac{1}{2}\log|\Sigma + \sigma_{\text{noise}}^2 \mathbf{I}| - \frac{n}{2}\log 2\pi \quad (1)$$

071

072  
073  
074  
075

where  $n$  is the number of data points. Both log-likelihood and predictive posterior can be computed efficiently in a Venture SP with an algorithm that resorts to Cholesky factorization(Rasmussen and Williams, 2006, chap. 2) resulting in a computational complexity of  $\mathcal{O}(n^3)$  in the number of data points.

076  
077  
078  
079  
080  
081  
082

The covariance function (or kernel) of a GP governs high-level properties of the observed data such as linearity, periodicity and smoothness. It comes with few free parameters that we call hyperparameters. Adjusting these results in minor changes, for example with regards to when two data points are treated similar. More drastically different covariance functions are achieved by changing the structure of the covariance function itself. Note that covariance function structures are compositional: adding or multiplying two valid covariance functions results in another valid covariance function.

083  
084  
085  
086

Venture includes the primitive `make_gp`, which takes as arguments a unary function `mean` and a binary (symmetric, positive-semidefinite) function `cov` and produces a function `g` distributed as a Gaussian process with the supplied mean and covariance. For example, a function  $g \sim \mathcal{GP}(0, \text{SE})$ , where `SE` is a squared-exponential covariance

$$087 \quad \text{SE}(x, x') = \sigma^2 \exp\left(\frac{(x - x')^2}{2\ell}\right)$$

088

089  
090  
091

with  $\sigma = 1$  and  $\ell = 1$ , can be instantiated as follows:

```
092 assume zero = make_const_func( 0.0 )
093 assume se = make_squaredexp( 1.0, 1.0 )
094 assume g = make_gp( zero, se )
```

095  
096  
097  
098  
099  
100

There are two ways two view `g` as a “random function.” In the first view, the `assume` directive that instantiates `g` does not use any randomness—only the subsequent calls to `g` do—and coherence constraints are upheld by the interpreter by keeping track of which evaluations of `g` exist in the current execution trace. Namely, if the current trace contains evaluations of `g` at the points  $x_1, \dots, x_N$  with return values  $y_1, \dots, y_N$ , then the next evaluation of `g` (say, jointly at the points  $x_{N+1}, \dots, x_{N+n}$ ) will be distributed according to the joint conditional distribution

$$101 \quad P((g x_{N+1}), \dots, (g x_{N+n}) | (g x_i) = y_i \text{ for } i = 1, \dots, N).$$

102

103  
104  
105  
106  
107

In the second view, `g` is a randomly chosen deterministic function, chosen from the space of all deterministic real-valued functions; in this view, the `assume` directive contains *all* the randomness, and subsequent invocations of `g` are deterministic. The first view is procedural and is faithful to the computation that occurs behind the scenes in Venture. The second view is declarative and is faithful to notations like “ $g \sim P(g)$ ” which are often used in mathematical treatments. Because a model program could make arbitrarily many calls to `g`, and the joint distribution on the return values of the

108 calls could have arbitrarily high entropy, it is not computationally possible in finite time to choose  
 109 the entire function  $g$  all at once as in the second view. Thus, it stands to reason that any computa-  
 110 tionally implementable notion of “nonparametric random functions” must involve incremental random  
 111 choices in one way or another, and Gaussian processes in Venture are no exception.  
 112

## 113 2.1 Memoization

114  
 115 Memoization is the practice of storing previously computed values of a function so that future calls  
 116 with the same inputs can be evaluated by lookup rather than recomputation. Although memoiza-  
 117 tion does not change the semantics of a deterministic program, it does change that of a stochastic  
 118 program (Goodman et al., 2008). The authors provide an intuitive example: let  $f$  be a function  
 119 that flips a coin and return “head” or “tails”. The probability that two calls of  $f$  are equivalent is  
 120 0.5. However, if the function call is memoized, it is 1. In fact, there is an infinite range of possible  
 121 caching policies (specifications of when to use a stored value and when to recompute), each poten-  
 122 tially having a different semantics. Any particular caching policy can be understood by random  
 123 world semantics (Poole, 1993; Sato, 1995) over the stochastic program: each possible world corre-  
 124 sponds to a mapping from function input sequence to function output sequence (McAllester et al.,  
 125 2008). In Venture, these possible worlds are first-class objects, and correspond to the *probabilistic*  
 126 *execution traces* (Mansinghka et al., 2014).

127 To transfer this idea to probabilistic programming, we now introduce a language construct called a  
 128 *statistical memoizer*. Suppose we have a function  $f$  which can be evaluated but we wish to learn  
 129 about the behavior of  $f$  using as few evaluations as possible. The statistical memoizer, which here  
 130 we give the name `gpmem`, was motivated by this purpose. It produces two outputs:  
 131

$$f \xrightarrow{\text{gpmem}} (f_{\text{probe}}, f_{\text{emu}}).$$

132 The function  $f_{\text{probe}}$  calls  $f$  and stores the output in a memo table, just as traditional memoization  
 133 does. The function  $f_{\text{emu}}$  is an online statistical emulator which uses the memo table as its training  
 134 data. A fully Bayesian emulator, modelling the true function  $f$  as a random function  $f \sim P(f)$ ,  
 135 would satisfy

$$(f_{\text{emu}} x_1 \dots x_k) \sim P(f(x_1), \dots, f(x_k) \mid f(x) = (f x) \text{ for each } x \text{ in memo table}).$$

136 Different implementations of the statistical memoizer can have different prior distributions  $P(f)$ ; in  
 137 this paper, we deploy a Gaussian process prior (implemented as `gpmem` below). Note that we require  
 138 the ability to sample  $f_{\text{emu}}$  jointly at multiple inputs because the values of  $f(x_1), \dots, f(x_k)$  will in  
 139 general be dependent. To illustrate the linguistic power of `gpmem` consider a simple model program  
 140 that uses `gpmem` procedure abstraction, and a loop:  
 141

```

1 define choose_next_point = proc(em) { gridsearch_argmax( em ) }
2
3 // Here stats( em ) is the memo table {(x_i, y_i)}.
4 // Take the (x,y) pair with the largest y.
5 define extract_answer = proc(em) { first(max( stats( em, second ))) }
6
7 // Hyper-parameter
8 assume theta = tag( quote( params ), 0, gamma(1,1))
9 assume (f_probe f_emu) = gpmem(f, make_squaredexp(1.0, theta))
10
11 for t...T:
12   f_probe( choose_next_point( f_emu ))
13   extract_answer( f_emu )
14
  
```

155 An equivalent model in typical statistics notation, written in a way that attempts to be as linguistically  
 156 faithful as possible, is

$$\begin{aligned}
 158 \quad & f^{(0)} \sim P(f) \\
 159 \quad & x^{(t)} = \arg \max_x f^{(t)}(x) \\
 160 \quad & f^{(t+1)} \sim P \left( f^{(t)} \mid f^{(t)}(x^{(t)}) = (f x^{(t)}) \right).
 \end{aligned}$$

162 The linguistic constructs for abstraction are not present in statistics notation. The equation  $x^{(t)} =$   
 163  $\arg \max_x f^{(t)}(x)$  has to be inlined in statistics notation: while something like  
 164

$$x^{(t)} \sim F(P^{(t)})$$

165 (where  $P^{(t)}$  is a probability measure, and  $F$  here plays the role of `choose_next_point`) is mathematically coherent, it is not what statisticians would ordinarily write. This lack of abstraction then  
 166 makes modularity, or the easy digestion of large but finely decomposable models, more difficult in  
 167 statistics notation.  
 168

169 Adding a single line to the program, such as  
 170

```
171     infer mh( quote( params ), one, 50 )
```

172 after line 12 to infer the parameters of the emulator, would require a significant refactoring and  
 173 elaboration of the statistics notation. One basic reason is that probabilistic programs are written pro-  
 174 cedurally, whereas statistical notation is declarative; reasoning declaratively about the dynamics of a  
 175 fundamentally procedural inference algorithm is often unwieldy due to the absence of programming  
 176 constructs such as loops and mutable state.  
 177

178 We implement `gpmem` by memoizing a target procedure in a wrapper that remembers previously  
 179 computed values. This comes with interesting implications: from the standpoint of computation, a  
 180 data set of the form  $\{(x_i, y_i)\}$  can be thought of as a function  $y = f_{\text{restr}}(x)$ , where  $f_{\text{restr}}$  is restricted  
 181 to only allow evaluation at a specific set of inputs  $x$  (Alg. 1). Modelling the data set with a GP then  
 182

---

#### Algorithm 1 Restricted Function

---

```
1: function  $f_{\text{RESTR}}(x)$ 
2:   if  $x \in \mathcal{D}$  then
3:     return  $\mathcal{D}[x]$ 
4:   else
5:     raise Exception('Illegal input')
6:   end if
7: end function
```

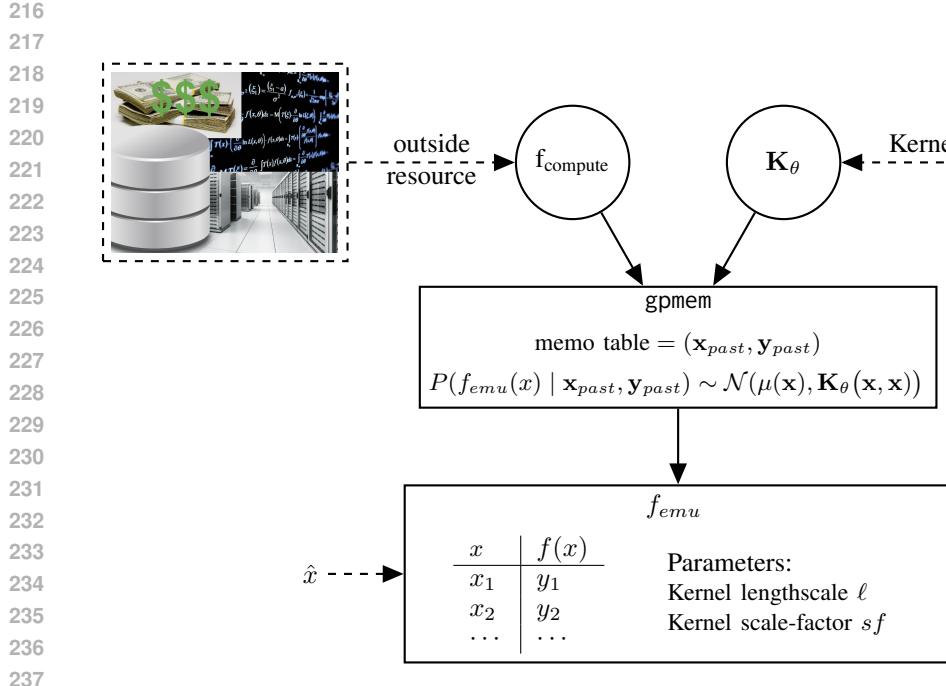
---

196  
 197 amounts to trying to learn a smooth function  $f_{\text{emu}}$  ("emu" stands for "emulator") which extends  $f$   
 198 to its full domain. Indeed, if  $f_{\text{restr}}$  is a foreign procedure made available as a black-box to Venture,  
 199 whose secret underlying pseudo code is in Alg. 1, then the `observe` code can be rewritten using  
 200 `gpmem` as in Listing 1 (where here the data set  $D$  has keys  $x[1], \dots, x[n]$ ):  
 201

Listing 1: Observation with `gpmem`

```
202
203
204 1  assume (f_compute f_emu) = gpmem( f_restr )
205 2  for i=1 to n:
206 3    predict f_compute( x[i] )
207 4    infer mh(quote(hyper-parameters), one, 100)
208 5    sample (f_emu( array( 1, 2, 3 ))
```

210  
 211 This rewriting has at least two benefits: (i) readability (in some cases), and (ii) amenability to active  
 212 learning. As to (i), the statistical code of creating a Gaussian process is replaced with a memoization-  
 213 like idiom, which will be more familiar to programmers. As to (ii), when using `gpmem`, it is quite  
 214 easy to decide incrementally which data point to sample next: for example, the loop from  $x[1]$  to  
 215  $x[n]$  could be replaced by a loop in which the next index  $i$  is chosen by a supplied decision rule. In  
 this way, we could use `gpmem` to perform online learning using only a subset of the available data.



```

239 define f = proc( x ) {
240     exp(-0.1*abs(x-2)) * 
241     10* cos(0.4*x) + 0.2
242 }
243 assume (f_compute f_emu) = gpmem( f, K_theta )
244 sample f_emu( array( -20, ..., 20))

245
246 predict f_compute( 12.6)
247
248 sample f_emu( array( -20, ..., 20))

249
250
251 predict f_compute( -6.4)
252
253 sample f_emu( array( -20, ..., 20))

254
255 observe f_emu( -3.1) = 2.60
256 observe f_emu( 7.8) = -7.60
257 observe f_emu( 0.0) = 10.19
258
259 sample f_emu( array( -20, ..., 20))

260
261
262 infer mh(quote(hyper-parameter), one, 50)
263
264 sample f_emu( array( -20, ..., 20))
265
266
267
268
269

```

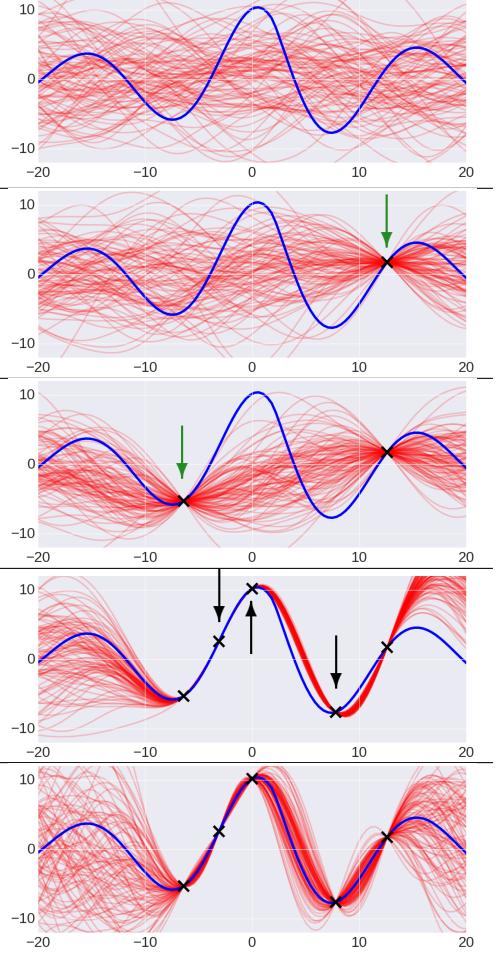



Figure 1: gpmem tutorial. The top shows a schematic of gpmem.  $f_{com} = f_{compute}$  probes an outside resource. This can be expensive (top left). Every probe is memoized and improves the GP-based emulator. Below the schematic we see a movie of the evolution of gpmem's state of believe of the world given certain Venture directives.



283  
284 Figure 2: (a) shows deterministic inference, that is optimization for an example of Neals synthetic  
285 function with an extreme outlier (green).  
286

### 287 3 Example Applications

288 gpmem is an elegant linguistic framework for function learning-related tasks. The technique allows  
289 language constructs from programming to help express models which would be cumbersome to  
290 express in statistics notation. We will now illustrate this with three example applications.  
291

#### 293 3.1 Hyperparameter estimation with structured hyper-prior

295 The probability of the hyper-parameters of a GP as defined above and given covariance function  
296 structure  $\mathbf{K}$  is:

$$297 \quad P(\boldsymbol{\theta} | \mathbf{D}, \mathbf{K}) = \frac{P(\mathbf{D} | \boldsymbol{\theta}, \mathbf{K})P(\boldsymbol{\theta} | \mathbf{K})}{P(\mathbf{D} | \mathbf{K})}. \quad (2)$$

300 Let the  $\mathbf{K}$  be the sum of a smoothing and a white noise (WN) kernel. For this case, Neal (1997)  
301 suggested the problem of outliers in data as a use-case for a hierarchical Bayesian treatment of  
302 Gaussian processes<sup>1</sup>. The work suggests a hierarchical system of hyper-parameterization. Here, we  
303 draw hyper-parameters from a  $\Gamma$  distributions:

$$304 \quad \ell^{(t)} \sim \Gamma(\alpha_1, \beta_1), \sigma^{(t)} \sim \Gamma(\alpha_2, \beta_2) \quad (3)$$

305 and in turn sample the  $\alpha$  and  $\beta$  from  $\Gamma$  distributions as well:

$$307 \quad \alpha_1^{(t)} \sim \Gamma(\alpha_\alpha^1, \beta_\alpha^1), \alpha_2^{(t)} \sim \Gamma(\alpha_\alpha^2, \beta_\alpha^2), \dots \quad (4)$$

309 One can represent this kind of model using gpmem (Fig. 3). Neal provides a custom inference  
310 algorithm setting and evaluates it using the following synthetic data problem. Let  $f$  be the underlying  
311 function that generates the data:

$$312 \quad f(x) = 0.3 + 0.4x + 0.5 \sin(2.7x) + \frac{1.1}{(1+x^2)} + \eta \quad \text{with } \eta \sim \mathcal{N}(0, \sigma) \quad (5)$$

315 We synthetically generate outliers by setting  $\sigma = 0.1$  in 95% of the cases and to  $\sigma = 1$  in the  
316 remaining cases. gpmem can capture the true underlying function within only 100 MH steps on the  
317 hyper-parameters to get a good approximation for their posterior (see Fig. 2). Note that Neal devises  
318 an additional noise model and performs a large number of Hybrid-Monte Carlo and Gibbs steps.

319 We illustrate the hyper-parameters by showing a comparison between  $\sigma$  (Fig. 2). We see that gpmem  
320 learns the posterior distribution well, the posterior even exhibits a bimodal histogram when sampling  
321  $\sigma$  100 times reflecting the two modes of data generation, that is normal noise and outliers.

322  
323 <sup>1</sup>In (Neal, 1997) the sum of an SE plus a constant kernel is used. We keep the WN kernel for illustrative  
324 purposes.

```

324
325
326
327
328
329
330
331
332
333 // Data and look-up function
334 define data = array(array(-1.87,0.13),..., array(1.67,0.81))
335 assume f_look_up = proc(index) {lookup( data, index)}
336
337
338 assume sf = tag(quote(hyper), 0, gamma(alpha_sf, beta_sf)))
339 assume l = tag(quote(hyper), 1, gamma(alpha_l, beta_l)))
340 assume sigma = tag(quote(hyper), 2, uniform_continuous(0, 2))
341
342 // The covariance function
343 assume se = make_squaredexp(sf, l)
344 assume wn = make_whitenoise(sigma)
345 assume composite_covariance = add_funcs(se, wn)
346
347 // Create a prober and emulator using gpmem
348 assume (f_compute, f_emu)
349 = gpmem(f_look_up, composite_covariance)
350
351 sample f_emu( array( -2, ..., 2))
352
353 // Observe all data points
354 for n ... N
355 observe f_emu(first(lookup(data,n)))
356 = second(lookup(data,n))
357 // Or: probe all data points
358 for n ... N
359 predict f_compute(first(lookup(data,n)))
360
361 sample f_emu( array( -2, ..., 2))
362
363 // Metropolis-Hastings
364 infer repeat( 100, do(
365 mh( quote(hyperhyper), one, 2),
366 mh( quote(hyper), one, 1)))
367
368 sample f_emu( array( -2, ..., 2))
369
370
371 // Optimization
372
373 infer map( quote(hyper), all, 0.01, 15)
374
375 sample f_emu( array( -2, ..., 2))
376
377
```

Hyper-Parameters for the Kernel:

---



---



---



---



---



---



---

Figure 3: Hyper-parameterization

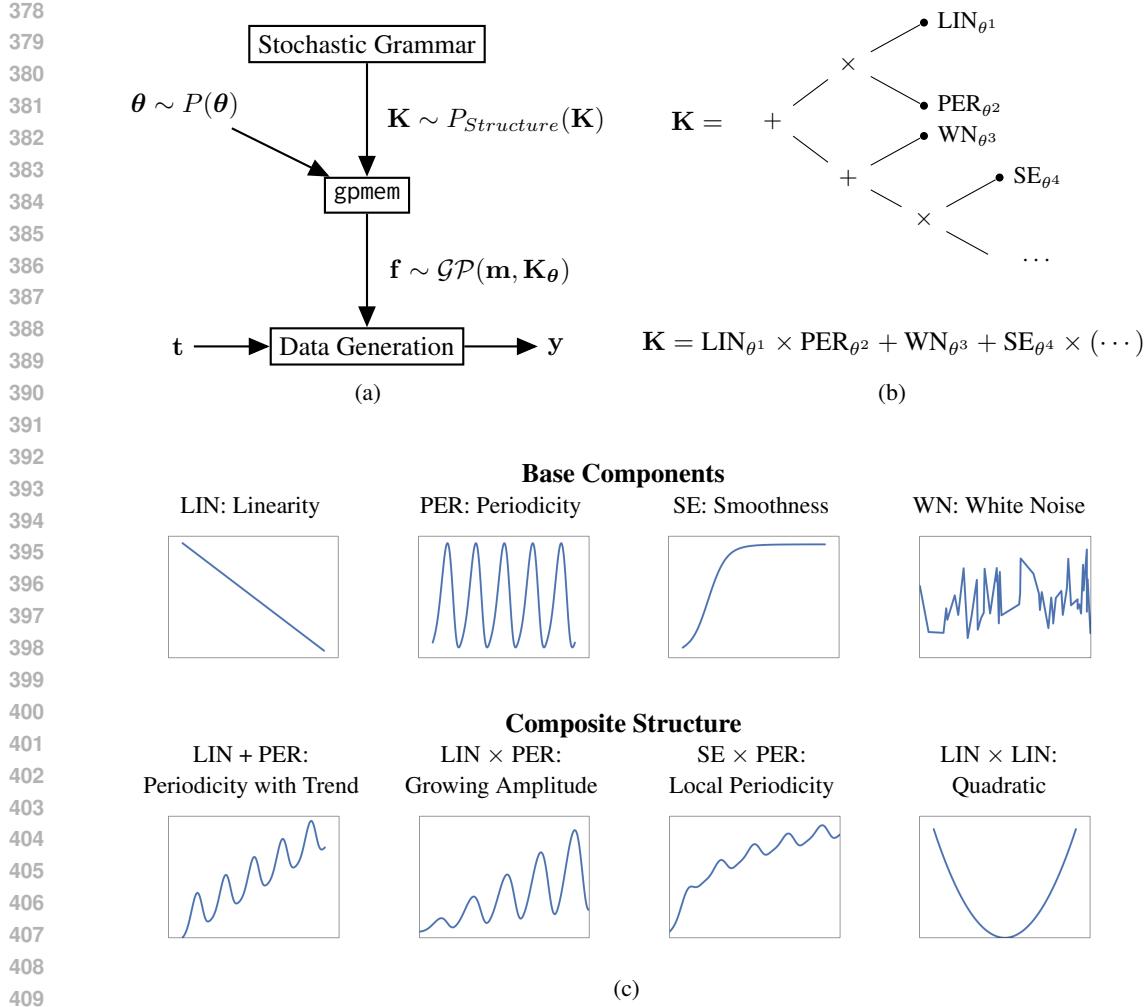


Figure 4: (a) Graphical description of Bayesian GP structure learning. (b) Composite structure. (c) The natural language interpretation of the structure.

### 3.2 Discovering symbolic models for time series

Inductive learning of symbolic expression for continuous-valued time series data is a hard task which has recently been tackled using a greedy search over the approximate posterior of the possible kernel compositions for GPs (Duvenaud et al., 2013; Lloyd et al., 2014)<sup>2</sup>.

With gpmem we can provide a fully Bayesian treatment of this, previously unavailable, using a stochastic grammar (see Fig. 4).

We deploy a probabilistic context free grammar for our prior on structures. An input of non-composite kernels (base kernels) is supplied to generate a posterior distributions of composite structure to express local and global aspects of the data.

We approximate the following intractable integrals of the expectation for the prediction:

$$\mathbb{E}[y^* | x^*, \mathbf{D}, \mathbf{K}] = \iint f(x^*, \theta, \mathbf{K}) P(\theta | \mathbf{D}, \mathbf{K}) P(\mathbf{K} | \Omega, s, n) d\theta d\mathbf{K}. \quad (6)$$

<sup>2</sup><http://www.automaticstatistician.com/>

432 This is done by sampling from the posterior probability distribution of the hyper-parameters and the  
 433 possible kernel:  
 434

$$435 \quad y^* \approx \frac{1}{T} \sum_{t=1}^T f(x^* | \boldsymbol{\theta}^{(t)}, \mathbf{K}^{(t)}). \quad (7)$$

437 In order to provide the sampling of the kernel, we introduce a stochastic process that simulates the  
 438 grammar for algebraic expressions of covariance function algebra:  
 439

$$440 \quad \mathbf{K}^{(t)} \sim P(\mathbf{K} | \boldsymbol{\Omega}, s, n) \quad (8)$$

442 Here, we start with the set of given base kernels and draw a random subset. For this subset of size  
 443  $n$ , we sample a set of possible operators  $\boldsymbol{\Omega}$  combining base kernels. The marginal probability of a  
 444 composite structure

$$445 \quad P(\mathbf{K} | \boldsymbol{\Omega}, s, n) = P(\boldsymbol{\Omega} | s, n) \times P(s | n) \times P(n), \quad (9)$$

446 is characterized by the prior  $P(n)$  on the number of base kernels used, the probability of a uniformly  
 447 chosen subset of the set of  $n$  possible covariance functions  
 448

$$449 \quad P(s | n) = \frac{n!}{|s|!}, \quad (10)$$

451 and the probability of sampling a global or a local structure, which is given by a binomial distribu-  
 452 tion:  
 453

$$454 \quad P(\boldsymbol{\Omega} | s, n) = \binom{n}{r} p_{+ \times}^k (1 - p_{+ \times})^{n-k}. \quad (11)$$

456 Many equivalent covariance structures can be sampled due to covariance function algebra and equiv-  
 457 alent representations with different parameterization (Lloyd et al., 2014). To inspect the posterior  
 458 of these equivalent structures we convert each kernel expression into a sum of products and subse-  
 459 quently simplify. All base kernels can be found in Appendix A, rules for this simplification can be  
 460 found in appendix B. The code for learning of kernel structure is as follows:

461  
 462  
 463  
 464  
 465  
 466  
 467  
 468  
 469  
 470  
 471  
 472  
 473  
 474  
 475  
 476  
 477  
 478  
 479  
 480  
 481  
 482  
 483  
 484  
 485

```

486
487 // GRAMMAR FOR KERNEL STRUCTURE
488 1 assume kernels = list(se, wn, lin, per, rq) // defined as above
489
490 // prior on the number of kernels
491 2 assume p_number_k = uniform_structure(n)
492 3 assume sub_s = tag(quote(grammar), 0,
493                      subset(kernels, p_number_k))
494
495 5 assume grammar = proc(l) {
496     // kernel composition
497     if (size(l) <= 1)
498         { first(l) }
499     else { if (bernoulli())
500           { add_funcs(first(l), grammar(rest(l))) }
501           else { mult_funcs(first(l), grammar(rest(l))) }
502     }
503
504 13 assume K = tag(quote(grammar), 1, grammar(sub_s))
505
506 // Probe all data points
507 15 for n ... N
508     predict f_compute(get_data_xs(n))
509
510 // PERFORMING INFERENCE
511 17 infer repeat(2000, do(
512     mh(quote(grammar), one, 1),
513     for kernel ∈ K
514         mh(quote(hyperkernel), one, 1)))
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539

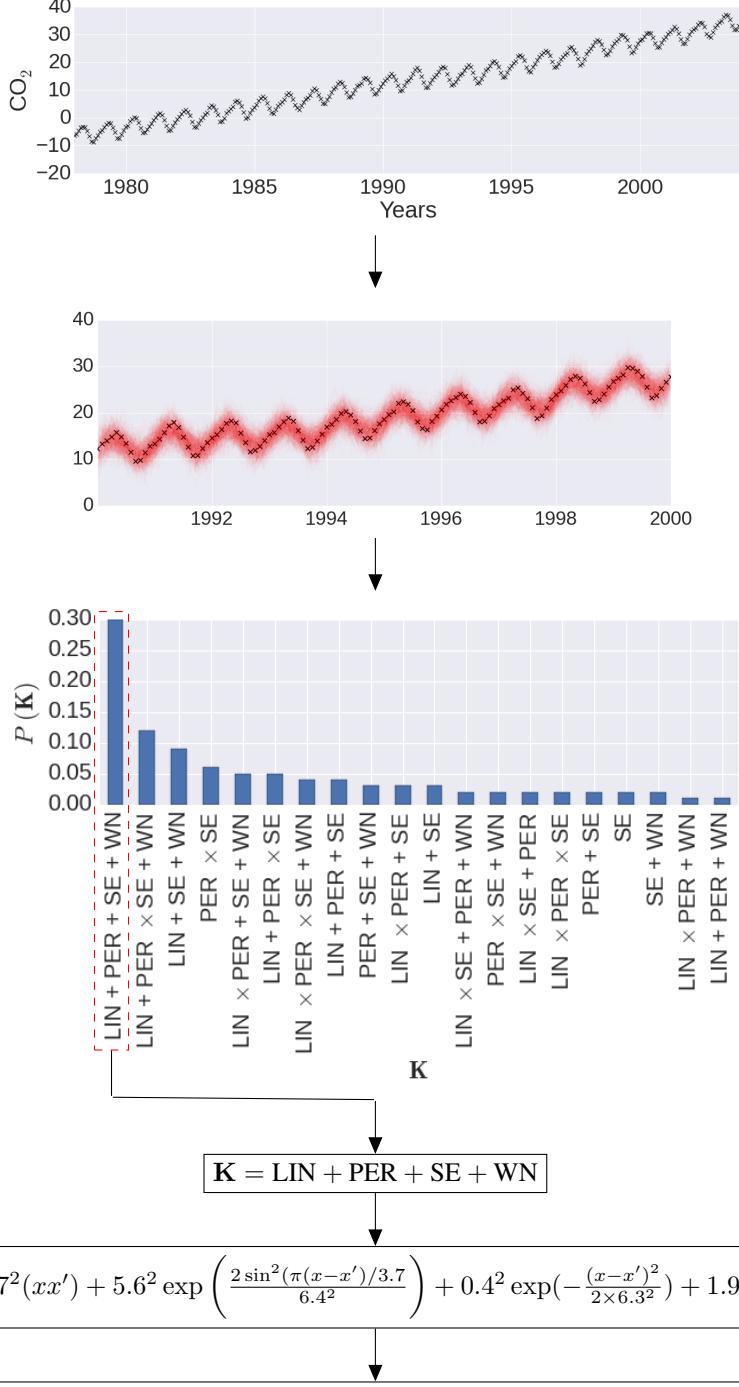
```

We defined the space of covariance structures in a way that allows us to produce results coherent with work presented in Automatic Statistician. For example, for the airline data set describing monthly totals of international airline passengers (Box et al., 1997, according to Duvenaud et al., 2013). Our most frequent sample is identical with the highest scoring result reported in previous work using a search-and-score method (Duvenaud et al., 2013) for the CO<sub>2</sub> data set (see Rasmussen and Williams, 2006 for a description) and the predictive capability is comparable. However, the components factor in a different way due to different parameterization of the individual base kernels. We see that the most probable alternatives for a structural description both recover the data dynamics (Fig. 12 for the airline data set).

Confident about our results, we can now query the data for certain structures being present. We illustrate this using the Mauna Loa data used in previous work on automated kernel discovery (Duvenaud et al., 2013). We assume a relatively simple hypothesis space consisting of only four kernels, a linear, a smoothing, a periodic and a white noise kernel. In this experiment, we resort to the white noise kernel instead RQ (similar to (Lloyd et al., 2014)). We can now run the algorithm, compute a posterior of structures (see Fig. 5). We can also query this posterior distribution for the marginal of certain simple structures to occur. We demonstrate this in Fig. 6

### 3.3 Bayesian optimization via Thompson sampling


In this section we present a probabilistic meta-program for Bayesian optimization, along with optimization routines implemented as concrete applications of the meta-program. Figure 7 shows a probabilistic meta-program for Bayesian optimization. We use the term “meta-program” because this program defines a generic rule for combining abstract components to perform optimization. The components are themselves subprograms, whose implementations may vary. A traditional example of generic programming is the tree search meta-program of Norvig (1992), written in LISP, which we present alongside the optimization meta-program for comparison.



**Qualitative Interpretation:** The the posterior peaks at a kernel structure with four additive components. Additive components hold globally, that is there are no higher level, qualitative aspects of the data that vary with the input space. The additive components are as follows: (i) a linearly increasing function or trend; (ii) a periodic function; (iii) a smooth function; and (iv) white noise.

Figure 5: Posterior of structure and qualitative, human interpretable reading. We take the raw data (top), compute a posterior distribution on structures (bar plot). We take the peak of this distribution ( $\text{LIN} + \text{PER} + \text{SE} + \text{WN}$ ) and show its human readable interpretation (left of bar plot). Below the bar plot show one sample of this structure with corresponding parameters. On the bottom, we sample from a GP with with this kernel-hyper-parameter combination.

$$P(\mathbf{K}) \approx \frac{1}{N} \sum_{n=1}^N f(\mathbf{K}_n) \text{ where } f(\mathbf{K}_n) = \begin{cases} 1, & \text{if } \mathbf{K}_n \in \mathbf{K}_{global}, \\ 0, & \text{otherwise.} \end{cases}$$



$$P(\mathbf{K}_a \wedge \mathbf{K}_b) \approx \frac{1}{N} \sum_{n=1}^N f(\mathbf{K}_n) \text{ where } f(\mathbf{K}_n) = \begin{cases} 1, & \text{if } \mathbf{K}_a \text{ and } \mathbf{K}_b \in_{global} \mathbf{K}_n, \\ 0, & \text{otherwise.} \end{cases}$$

$$P(\mathbf{K}_a \vee \mathbf{K}_b) = P(\mathbf{K}_a) + P(\mathbf{K}_b) - P(\mathbf{K}_a \wedge \mathbf{K}_b)$$


What is the probability of a trend, a recurring pattern **and** noise in the data?

$$P((\text{LIN} \vee \text{LIN} \times \text{SE}) \wedge (\text{PER} \vee \text{PER} \times \text{SE} \vee \text{PER} \times \text{LIN}) \wedge (\text{WN} \vee \text{LIN} \times \text{WN})) = 0.36$$

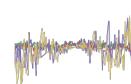
**Is there a trend?**  
 $P(\text{LIN} \vee \text{LIN} \times \text{SE}) = 0.65$




A linear trend?  $P(\text{LIN}) = 0.63$  A smooth trend?  $P(\text{LIN} \times \text{SE}) = 0.02$




**Is there noise?**



Heteroskedastic noise?    White noise?  
 $P(\text{LIN} \times \text{WN}) = 0$      $P(\text{WN}) = 0.75$



Is there repeating structure?  
 $P(\text{PER} \vee \text{PER} \times \text{SE} \vee \text{PER} \times \text{LIN}) = 0.73$



PER = 0.32



PER X



$\times \text{LJN} = 0.07$



Figure 6: We can query the data if some logical statements are probable to be true, for example, is it true that there is a trend?

```

648
649
650
651
652
653
654
655
656
657
658
659
660
661
662 (defun tree-search (states goal-p successors combiner)
663   "Find a state that satisfies goal-p. Start with states,
664   and search according to successors and combiner."
665   (cond ((null states) fail)
666         ((funcall goal-p (first states)) (first states))
667         (t (tree-search
668             (funcall combiner
669                 (funcall successors (first states))
670                 (rest states))
671                 goal-p successors combiner))))
```

```

672
673
674 (define optimize
675   (lambda (probe do-search search-state-box
676             post-probe-inference extract-answer finished?)
677     (let loop ()
678       (if (finished?)
679           (extract-answer)
680           (begin
681             (do-search)
682             (if (contents-changed? search-state-box)
683                 (predict (probe ,(contents search-state-box))))
684                 (post-probe-inference)
685                 (loop)))))))
```

Figure 7: Top: Norvig's LISP meta-program for tree search. Bottom: A probabilistic meta-program for Bayesian optimization, written in VentureScheme, an abstract, Scheme-like syntax that can be used for Venture.

```

690
691
692
693
694
695
696
697
698
699
700
701
```

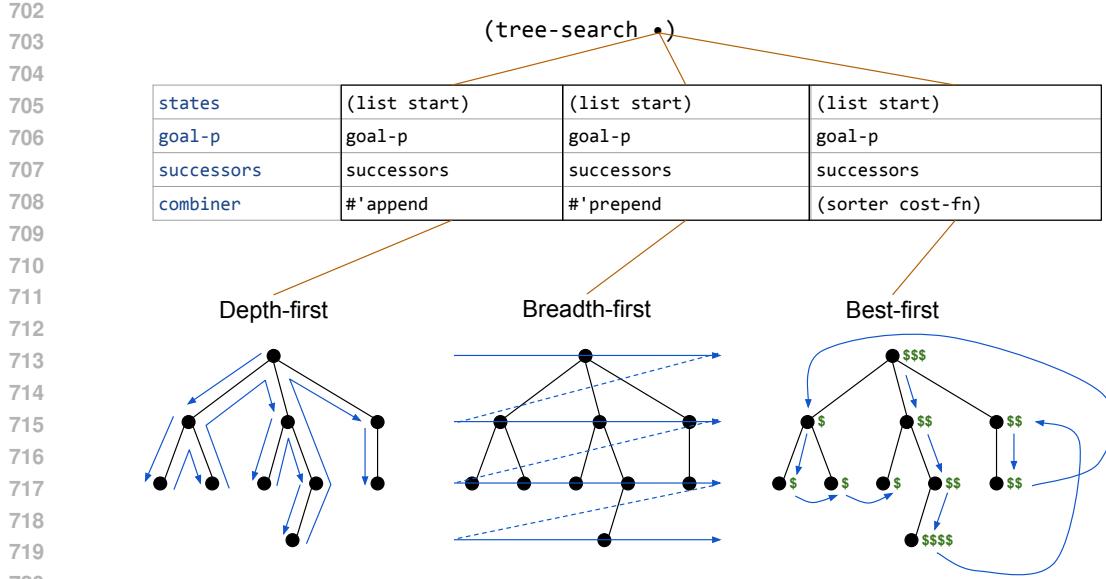


Figure 8: Invocations of Norvig’s tree search meta-procedure with different arguments to perform a variety of different tree searches.

Figure 8 shows the flexibility of Norvig’s meta-program. The procedure `tree-search` can be turned into a concrete implementation of depth-first, breadth-first or best-first search (among others) by simply supplying the appropriate combiner and set of initial states. Similarly, the meta-procedure `optimize` can make use of a variety of statistical emulators (e.g., GP-based emulators or polynomial regressions), a variety of acquisition heuristics (e.g., randomized grid search or Gaussian drift) and even a variety of rules for choosing an optimum based on the final state of the program (e.g., return the best probe so far, or return the analytically determined optimum of the emulator).

The specifications for the arguments to the meta-procedure `optimize` in Figures 7 are as follows:

- **probe:** A procedure for querying the true function  $f$  (and storing its value in the appropriate table, if caching is desired).
- **do\_search:** A procedure to search for a new probe point and, if one is found, store it in `search_state_box`.
- **search\_state\_box:** A container for the next value to be probed. In each iteration of the loop, if the contents of `search_state_box` have changed, a new probe is performed.
- **post\_probe\_inference:** Any inference instructions which should be run after each probe, such as inferring hyperparameters.
- **extract\_answer:** A procedure to produce an approximate optimum, based on the probes and inference that have been done so far.
- **finished?:** A procedure to decide whether optimization should be truncated here or should continue.

The meta-program is quite simple: Until `finished?` decides that optimization is finished, we choose new probe points and call `probe` on them. After each probe, we perform post-probe inference. Once optimization is finished, we call `extract_answer` to obtain an approximate optimum.

### Thompson Sampling

We introduce Thompson sampling, the basic solution strategy underlying the Bayesian optimization with `gpmem`. Thompson sampling (Thompson 1933) is a widely used Bayesian framework for addressing the trade-off between exploration and exploitation in multi-armed (or continuum-armed) bandit problems. We cast the multi-armed bandit problem as a one-state Markov decision process,

756 and describe how Thompson sampling can be used to choose actions for that Markov Decision  
757 Processes (MDP).  
758

759 Here, an agent is to take a sequence of actions  $a_1, a_2, \dots$  from a (possibly infinite) set of possible  
760 actions  $\mathcal{A}$ . After each action, a reward  $r \in \mathbb{R}$  is received, according to an unknown conditional  
761 distribution  $P_{\text{true}}(r | a)$ . The agent's goal is to maximize the total reward received for all actions in  
762 an online manner. In Thompson sampling, the agent accomplishes this by placing a prior distribution  
763  $P(\vartheta)$  on the possible "contexts"  $\vartheta \in \Theta$ . Here a context is a believed model of the conditional  
764 distributions  $\{P(r | a)\}_{a \in \mathcal{A}}$ , or at least, a believed statistic of these conditional distributions which  
765 is sufficient for deciding an action  $a$ . If actions are chosen so as to maximize expected reward, then  
766 one such sufficient statistic is the believed conditional mean  $V(a | \vartheta) = \mathbb{E}[r | a; \vartheta]$ , which can be  
767 viewed as a believed value function. For consistency with what follows, we will assume our context  
768  $\vartheta$  takes the form  $(\theta, \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}})$  where  $\mathbf{a}_{\text{past}}$  is the vector of past actions,  $\mathbf{r}_{\text{past}}$  is the vector of their  
769 rewards, and  $\theta$  (the "semicontext") contains any other information that is included in the context.  
770

In this setup, Thompson sampling has the following steps:

---

**Algorithm 2** Thompson sampling.

---

771 Repeat as long as desired:  
772

- 773 1. **Sample.** Sample a semicontext  $\theta \sim P(\theta)$ .
  - 774 2. **Search (and act).** Choose an action  $a \in \mathcal{A}$  which (approximately) maximizes  $V(a | \vartheta) =$   
775  $\mathbb{E}[r | a; \vartheta] = \mathbb{E}[r | a; \theta, \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}}]$ .
  - 776 3. **Update.** Let  $r_{\text{true}}$  be the reward received for action  $a$ . Update the believed distribution on  
777  $\theta$ , i.e.,  $P(\theta) \leftarrow P_{\text{new}}(\theta)$  where  $P_{\text{new}}(\theta) = P(\theta | a \mapsto r_{\text{true}})$ .
- 

781  
782 Note that when  $\mathbb{E}[r | a; \vartheta]$  (under the sampled value of  $\theta$  for some points  $a$ ) is far from the true value  
783  $\mathbb{E}_{P_{\text{true}}}[r | a]$ , the chosen action  $a$  may be far from optimal, but the information gained by probing  
784 action  $a$  will improve the belief  $\vartheta$ . This amounts to "exploration." When  $\mathbb{E}[r | a; \vartheta]$  is close to the  
785 true value except at points  $a$  for which  $\mathbb{E}[r | a; \vartheta]$  is low, exploration will be less likely to occur, but  
786 the chosen actions  $a$  will tend to receive high rewards. This amounts to "exploitation." The trade-off  
787 between exploration and exploitation is illustrated in Figure 9. Roughly speaking, exploration will  
788 happen until the context  $\vartheta$  is reasonably sure that the unexplored actions are probably not optimal,  
789 at which time the Thompson sampler will exploit by choosing actions in regions it knows to have  
790 high value.

791 Typically, when Thompson sampling is implemented, the search over contexts  $\vartheta \in \Theta$  is limited  
792 by the choice of representation. In traditional programming environments,  $\theta$  often consists of a few  
793 numerical parameters for a family of distributions of a fixed functional form. With work, a mixture of  
794 a few functional forms is possible; but without probabilistic programming machinery, implementing  
795 a rich context space  $\Theta$  would be an unworkably large technical burden. In a probabilistic programming  
796 language, however, the representation of heterogeneously structured or infinite-dimensional context  
797 spaces is quite natural. Any computable model of the conditional distributions  $\{P(r | a)\}_{a \in \mathcal{A}}$  can  
798 be represented as a stochastic procedure  $(\lambda(a) \dots)$ . Thus, for computational Thompson sampling,  
799 the most general context space  $\hat{\Theta}$  is the space of program texts. Any other context space  $\Theta$  has a  
800 natural embedding as a subset of  $\hat{\Theta}$ .  
801

## 802 A Mathematical Specification

803 Our mathematical specification assert the following properties:  
804

- 805 • The regression function has a Gaussian process prior.
- 806 • The actions  $a_1, a_2, \dots \in \mathcal{A}$  are chosen by a Metropolis-like search strategy with Gaussian  
807 drift proposals.
- 808 • The hyperparameters of the Gaussian process are inferred using Metropolis–Hastings sam-  
809 pling after each action.

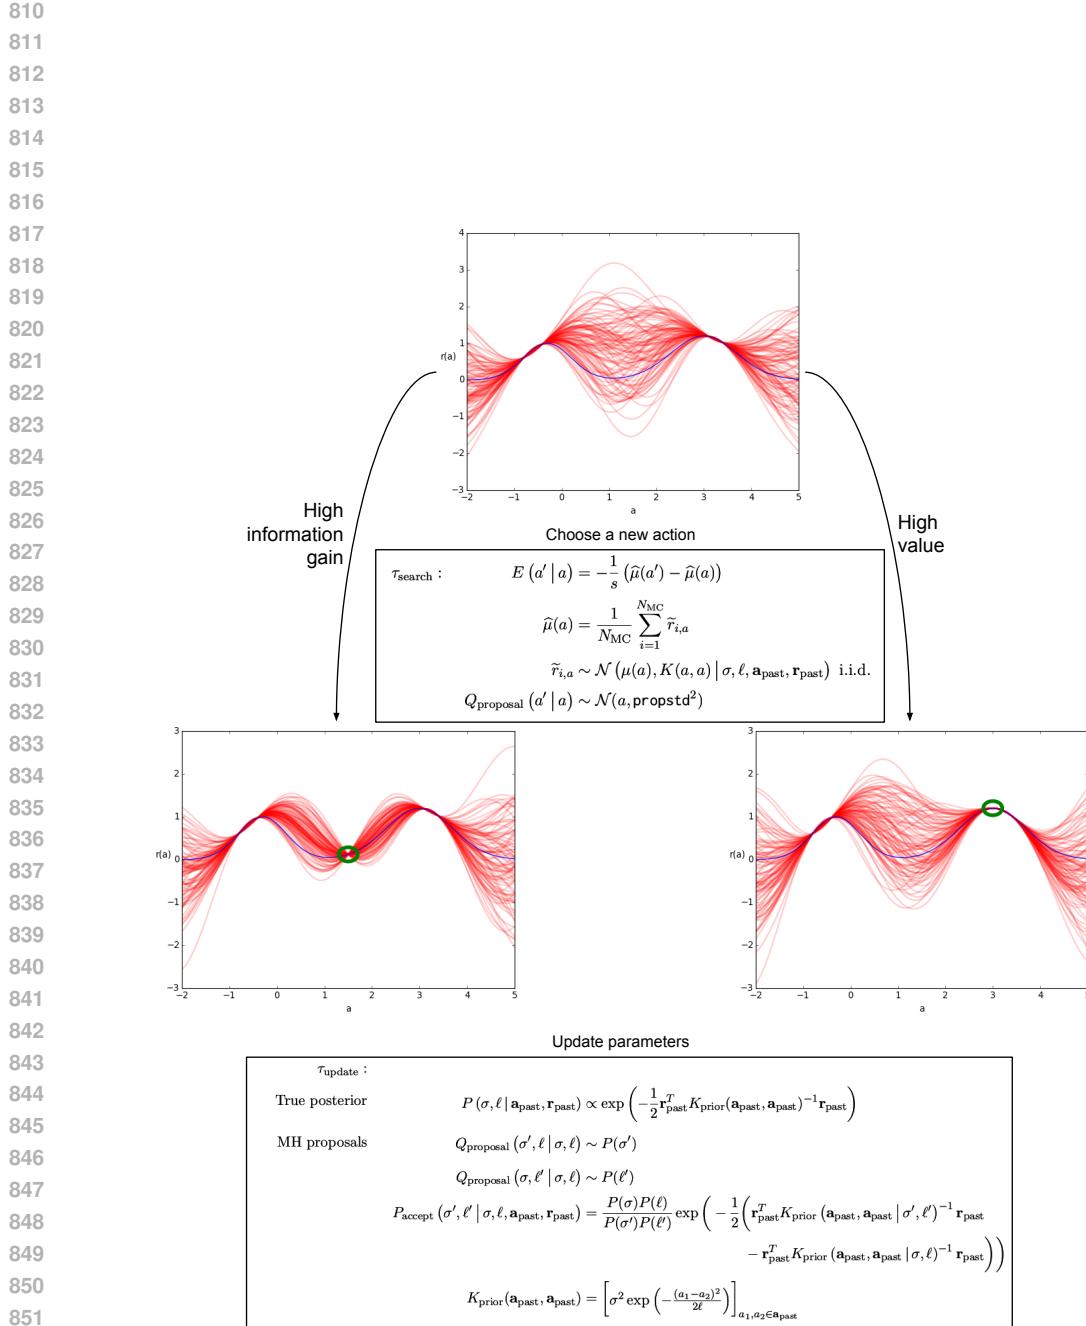


Figure 9: Two possible actions (in green) for an iteration of Thompson sampling. The believed distribution on the value function  $V$  is depicted in red. In this example, the true reward function is deterministic, and is drawn in blue. The action on the right receives a high reward, while the action on the left receives a low reward but greatly improves the accuracy of the believed distribution on  $V$ . The transition operators  $\tau_{\text{search}}$  and  $\tau_{\text{update}}$  are described in Section 3.3.

In this version of Thompson sampling, the contexts  $\vartheta$  are Gaussian processes over the action space  $\mathcal{A} = [-20, 20] \subseteq \mathbb{R}$ . That is,

$$V \sim \mathcal{GP}(\mu, K),$$

where the mean  $\mu$  is a computable function  $\mathcal{A} \rightarrow \mathbb{R}$  and the covariance  $K$  is a computable (symmetric, positive-semidefinite) function  $\mathcal{A} \times \mathcal{A} \rightarrow \mathbb{R}$ . This represents a Gaussian process  $\{R_a\}_{a \in \mathcal{A}}$ , where  $R_a$  represents the reward for action  $a$ . Computationally, we represent a context as a data structure

$$\vartheta = (\theta, \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}}) = (\mu_{\text{prior}}, K_{\text{prior}}, \eta, \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}}),$$

where  $\mu_{\text{prior}}$  is a procedure to be used as the prior mean function, w.l.o.g. we set  $\mu_{\text{prior}} \equiv 0$ .  $K_{\text{prior}}$  is a procedure to be used as the prior covariance function, parameterized by  $\eta$ .

The posterior mean and covariance for such a context  $\vartheta$  are gotten by the usual conditioning formulas (assuming, for ease of exposition as above, that the prior mean is zero):<sup>3</sup>

$$\begin{aligned}\mu(\mathbf{a}) &= \mu(\mathbf{a} | \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}}) \\ &= K_{\text{prior}}(\mathbf{a}, \mathbf{a}_{\text{past}}) K_{\text{prior}}(\mathbf{a}_{\text{past}}, \mathbf{a}_{\text{past}})^{-1} \mathbf{r}_{\text{past}} \\ K(\mathbf{a}, \mathbf{a}) &= K(\mathbf{a}, \mathbf{a} | \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}}) \\ &= K_{\text{prior}}(\mathbf{a}, \mathbf{a}) - K_{\text{prior}}(\mathbf{a}, \mathbf{a}_{\text{past}}) K_{\text{prior}}(\mathbf{a}_{\text{past}}, \mathbf{a}_{\text{past}})^{-1} K_{\text{prior}}(\mathbf{a}_{\text{past}}, \mathbf{a}).\end{aligned}$$

Note that the context space  $\Theta$  is not a finite-dimensional parametric family, since the vectors  $\mathbf{a}_{\text{past}}$  and  $\mathbf{r}_{\text{past}}$  grow as more samples are taken.  $\Theta$  is, however, representable as a computational procedure together with parameters and past samples, as we do in the representation  $\vartheta = (\mu_{\text{prior}}, K_{\text{prior}}, \eta, \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}})$ .

We combine the Update and Sample steps of Algorithm 2 by running a Metropolis–Hastings (MH) sampler whose stationary distribution is the posterior  $P(\theta | \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}})$ . The functional forms of  $\mu_{\text{prior}}$  and  $K_{\text{prior}}$  are fixed in our case, so inference is only done over the parameters  $\eta = \{\sigma, \ell\}$ ; hence we equivalently write  $P(\sigma, \ell | \mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}})$  for the stationary distribution. We make MH proposals to one variable at a time, using the prior as proposal distribution:

$$Q_{\text{proposal}}(\sigma', \ell' | \sigma, \ell) = P(\sigma')$$

and

$$Q_{\text{proposal}}(\sigma, \ell' | \sigma, \ell) = P(\ell').$$

The MH acceptance probability for such a proposal is

$$P_{\text{accept}}(\sigma', \ell' | \sigma, \ell) = \min \left\{ 1, \frac{Q_{\text{proposal}}(\sigma, \ell | \sigma', \ell')}{Q_{\text{proposal}}(\sigma', \ell' | \sigma, \ell)} \cdot \frac{P(\mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}} | \sigma', \ell')}{P(\mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}} | \sigma, \ell)} \right\}$$

Because the priors on  $\sigma$  and  $\ell$  are uniform in our case, the term involving  $Q_{\text{proposal}}$  equals 1 and we have simply

$$\begin{aligned}P_{\text{accept}}(\sigma', \ell' | \sigma, \ell) &= \min \left\{ 1, \frac{P(\mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}} | \sigma', \ell')}{P(\mathbf{a}_{\text{past}}, \mathbf{r}_{\text{past}} | \sigma, \ell)} \right\} \\ &= \min \left\{ 1, \exp \left( -\frac{1}{2} \left( \mathbf{r}_{\text{past}}^T K_{\text{prior}}(\mathbf{a}_{\text{past}}, \mathbf{a}_{\text{past}} | \sigma', \ell')^{-1} \mathbf{r}_{\text{past}} \right. \right. \right. \\ &\quad \left. \left. \left. - \mathbf{r}_{\text{past}}^T K_{\text{prior}}(\mathbf{a}_{\text{past}}, \mathbf{a}_{\text{past}} | \sigma, \ell)^{-1} \mathbf{r}_{\text{past}} \right) \right) \right\}.\end{aligned}$$

The proposal and acceptance/rejection process described above define a transition operator  $\tau_{\text{update}}$  which is iterated a specified number of times; the resulting state of the MH Markov chain is taken as the sampled semicontext  $\theta$  in Step 1 of Algorithm 2.

For Step 2 (Search) of Thompson sampling, we explore the action space using an MH-like transition operator  $\tau_{\text{search}}$ . As in MH, each iteration of  $\tau_{\text{search}}$  produces a proposal which is either accepted or rejected, and the state of this Markov chain after a specified number of steps is the new action  $a$ .

---

<sup>3</sup>Here, for vectors  $\mathbf{a} = (a_i)_{i=1}^n$  and  $\mathbf{a}' = (a'_i)_{i=1}^{n'}$ ,  $\mu(\mathbf{a})$  denotes the vector  $(\mu(a_i))_{i=1}^n$  and  $K(\mathbf{a}, \mathbf{a}')$  denotes the matrix  $[K(a_i, a'_j)]_{1 \leq i \leq n, 1 \leq j \leq n'}$ .

918 The Markov chain's initial state is the most recent action, and the proposal distribution is Gaussian  
 919 drift:

$$Q_{\text{proposal}}(a' | a) \sim \mathcal{N}(a, \text{propstd}^2),$$

920 where the drift width `propstd` is specified ahead of time. The acceptance probability of such a  
 921 proposal is

$$P_{\text{accept}}(a' | a) = \min \{1, \exp(-E(a' | a))\},$$

922 where the energy function  $E(\bullet | a)$  is given by a Monte Carlo estimate of the difference in value  
 923 from the current action:

$$E(a' | a) = -\frac{1}{s} (\hat{\mu}(a') - \hat{\mu}(a))$$

924 where

$$\hat{\mu}(a) = \frac{1}{N_{\text{avg}}} \sum_{i=1}^{N_{\text{avg}}} \tilde{r}_{i,a}$$

925 and

$$\tilde{r}_{i,a} \sim \mathcal{N}(\mu(a), K(a, a))$$

926 and  $\{\tilde{r}_{i,a}\}_{i=1}^{N_{\text{avg}}}$  are i.i.d. for a fixed  $a$ . Here the temperature parameter  $s \geq 0$  and the population size  
 927  $N_{\text{avg}}$  are specified ahead of time. Proposals of estimated value higher than that of the current action  
 928 are always accepted, while proposals of estimated value lower than that of the current action are  
 929 accepted with a probability that decays exponentially with respect to the difference in value. The  
 930 rate of the decay is determined by the temperature parameter  $s$ , where high temperature corresponds  
 931 to generous acceptance probabilities. For  $s = 0$ , all proposals of lower value are rejected; for  
 932  $s = \infty$ , all proposals are accepted. For points  $a$  at which the posterior mean  $\mu(a)$  is low but the  
 933 posterior variance  $K(a, a)$  is high, it is possible (especially when  $N_{\text{avg}}$  is small) to draw a “wild”  
 934 value of  $\hat{\mu}(a)$ , resulting in a favorable acceptance probability.

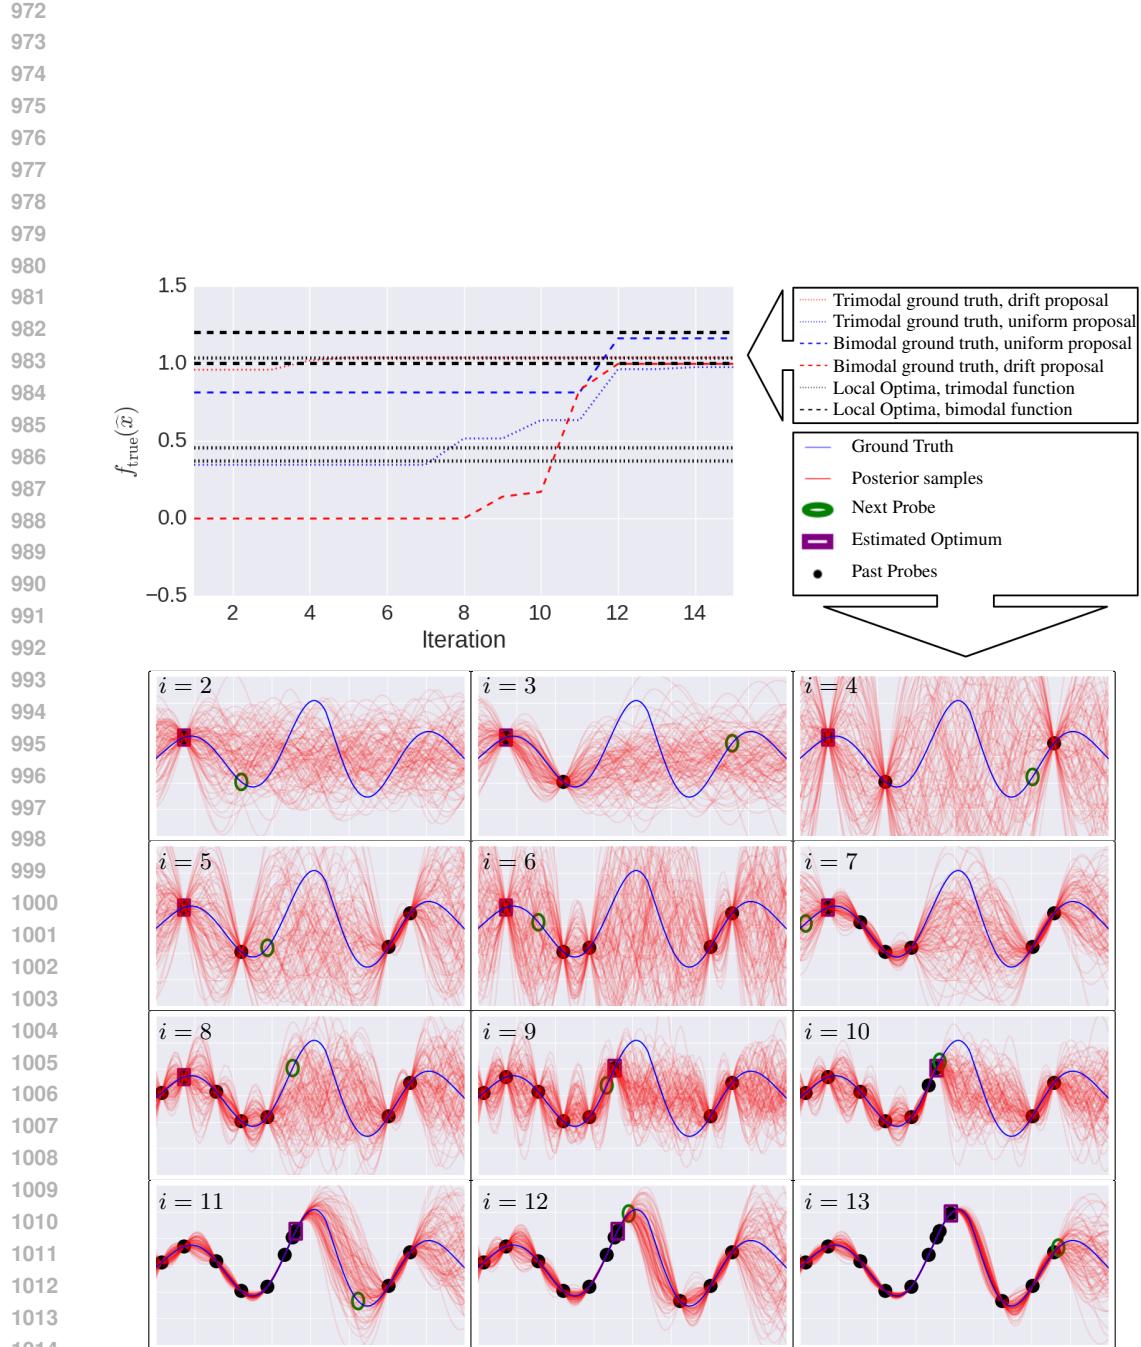
935 Indeed, taking an action  $a$  with low estimated value but high uncertainty serves the useful function  
 936 of improving the accuracy of the estimated value function at points near  $a$  (see Figure 9).<sup>4,5</sup> We see  
 937 a complete probabilistic program with `gpmem` implementing Bayesian optimization with Thompson  
 938 Sampling below (Listing 2).

939  
 940  
 941  
 942  
 943  
 944  
 945  
 946  
 947  
 948  
 949  
 950  
 951  
 952  
 953  
 954  
 955  
 956  
 957  
 958  
 959  
 960  
 961  
 962  
 963  
 964  
 965  
 966  
 967  
 968  
 969

---

<sup>4</sup>At least, this is true when we use a smoothing prior covariance function such as the squared exponential.

<sup>5</sup>For this reason, we consider the sensitivity of  $\hat{\mu}$  to uncertainty to be a desirable property; indeed, this is why we use  $\hat{\mu}$  rather than the exact posterior mean  $\mu$ .



1015 Figure 10: Top: the estimated optimum over time. Blue and Red represent optimization with uniform  
1016 and Gaussian drift proposals. Black lines indicate the local optima of the true functions. Bottom:  
1017 a sequence of actions. Depicted are iterations 7-12 with uniform proposals.

1018  
1019  
1020  
1021  
1022  
1023  
1024  
1025

```

1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1040
1041
1042
1043
1044
1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079

    Listing 2: Bayesian optimization using gpmem

1 assume sf = tag(quote(hyper), 0, uniform_continuous(0, 10))
2 assume l = tag(quote(hyper), 1, uniform_continuous(0, 10))
3 assume se = make_squaredexp(sf, 1)
4 assume blackbox_f = get_bayesopt_blackbox()
5 assume (f_compute, f_emulate) = gpmem(blackbox_f, se)

6 // A naive estimate of the argmax of the given function
7 define mc_argmax = proc(func) {
8     candidate_xs = mapv(proc(i) {uniform_continuous(-20, 20)},
9                          arange(20));
10    candidate_ys = mapv(func, candidate_xs);
11    lookup(candidate_xs, argmax_of_array(candidate_ys))
12};

13 // Shortcut to sample the emulator at a single point without packing
14 // and unpacking arrays
15 define emulate_pointwise = proc(x) {
16     run(sample(lookup(f_emulate(array(unquote(x))), 0)))
17};

18 // Main inference loop
19 infer repeat(15, do(pass,
20     // Probe V at the point mc_argmax(emulate_pointwise)
21     predict(f_compute(unquote(mc_argmax(emulate_pointwise)))),
22     // Infer hyperparameters
23     mh(quote(hyper), one, 50)));

```

## 4 Discussion

We provided gpmem, an elegant linguistic framework for function learning-related tasks such as Bayesian optimization and GP kernel structure learning. We highlighted how gpmem overcomes shortcomings of the notations currently used in statistics, and how language constructs from programming allow the expression of models which would be cumbersome (prohibitively so, in some cases) to express in statistics notation. We evaluated our contribution on a range of hard problems for state-of-the-art Bayesian nonparametrics.

## Appendix

## A Covariance Functions

$$\text{SE} = \sigma^2 \exp\left(-\frac{(x - x')^2}{2\ell^2}\right) \quad (12)$$

$$\text{LIN} = \sigma^2(xx') \quad (13)$$

$$C = \sigma^2 \quad (14)$$

$$\text{WN} = \sigma^2 \delta_{x,x'} \quad (15)$$

$$RQ = \sigma^2 \left( 1 + \frac{(x - x')^2}{2\alpha\ell^2} \right)^{-\alpha} \quad (16)$$

$$\text{PER} = \sigma^2 \exp\left(\frac{2 \sin^2(\pi(x - x')/p)}{\ell^2}\right). \quad (17)$$

## B Covariance Simplification

$SE \times SE$	$\rightarrow SE$
$\{SE, PER, C, WN\} \times WN$	$\rightarrow WN$
$LIN + LIN$	$\rightarrow LIN$
$\{SE, PER, C, WN, LIN\} \times C$	$\rightarrow \{SE, PER, C, WN, LIN\}$

Rule 1 is derived as follows:

$$\begin{aligned}
\sigma_c^2 \exp\left(-\frac{(x - x')^2}{2\ell_c^2}\right) &= \sigma_a^2 \exp\left(-\frac{(x - x')^2}{2\ell_a^2}\right) \times \sigma_b^2 \exp\left(-\frac{(x - x')^2}{2\ell_b^2}\right) \\
&= \sigma_c^2 \exp\left(-\frac{(x - x')^2}{2\ell_a^2}\right) \times \exp\left(-\frac{(x - x')^2}{2\ell_b^2}\right) \\
&= \sigma_c^2 \exp\left(-\frac{(x - x')^2}{2\ell_a^2} - \frac{(x - x')^2}{2\ell_b^2}\right) \\
&= \sigma_c^2 \exp\left(-\frac{(x - x')^2}{2\ell_c^2}\right)
\end{aligned} \tag{18}$$

For stationary kernels that only depend on the lag vector between  $x$  and  $x'$  it holds that multiplying such a kernel with a WN kernel we get another WN kernel (Rule 2). Take for example the SE kernel:

$$\sigma_a^2 \exp\left(-\frac{(x-x')^2}{2\ell_a^2}\right) \times \sigma_b \delta_{x,x'} = \sigma_a \sigma_b \delta_{x,x'} \quad (19)$$

Rule 3 is derived as follows:

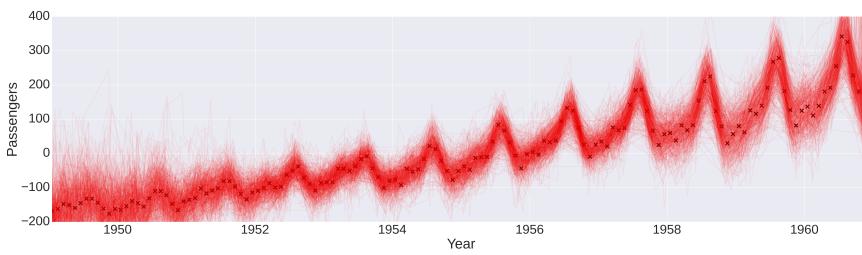
$$\theta_c(x \times x') = \theta_a(x \times x') + \theta_b(x \times x') \quad (20)$$

Multiplying any kernel with a constant obviously changes only the scale parameter of a kernel (Rule 4).

## C Additional Structure Learning Results

Below we depict additional results for the airline data set using a hypothesis space of LIN, PER, SE and RQ covariance functions as base kernels for the composition. Fig 11 shows the training data and posterior samples drawn from the GP with posterior composite structure.

Fig. 12 shows why it advantageous to be Bayesian here. We see two possible hypotheses for the composite structure of  $\mathbf{K}$ .



1134  
1135  
1136  
1137  
1138  
1139  
1140  
1141  
1142  
1143  
1144  
1145  
1146  
1147  
1148  
1149  
1150  
1151  
1152  
1153  
1154  
1155  
1156  
1157  
1158  
1159  
1160  
1161  
1162  
1163  
1164  
1165  
1166  
1167  
1168  
1169  
1170  
1171  
1172  
1173  
1174  
1175  
1176  
1177  
1178  
1179  
1180  
1181  
1182  
1183  
1184  
1185  
1186  
1187  
Figure 11: Data (black x) and posterior samples (red) for the airline data set.



1150  
1151  
1152  
1153  
1154  
1155  
1156  
1157  
1158  
1159  
1160  
1161  
1162  
1163  
1164  
1165  
1166  
1167  
1168  
1169  
1170  
1171  
1172  
1173  
1174  
1175  
1176  
1177  
1178  
1179  
1180  
1181  
1182  
1183  
1184  
1185  
1186  
1187  
Figure 12: We see two possible hypotheses for the composite structure of  $\mathbf{K}$ . (a) Most frequent sample drawn from the posterior on structure. We have found two global components. First, a smooth trend ( $\text{LIN} \times \text{SE}$ ) with a non-linear increasing slope. Second, a periodic component with increasing variation and noise. (b) Second most frequent sample drawn from the posterior on structure. We found one global component. It is comprised of local changes that are periodic and with changing variation.

## References

- Barry, D. (1986). Nonparametric bayesian regression. *The Annals of Statistics*, 14(3):934–953.
- Box, G. E. P., Jenkins, G. M., and Reinsel, G. C. (1997). *Time series analysis: forecasting and control*.
- Damianou, A. and Lawrence, N. (2013). Deep gaussian processes. In *Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 207–215.
- Duvenaud, D., Lloyd, J. R., Grosse, R., Tenenbaum, J., and Ghahramani, Z. (2013). Structure discovery in nonparametric regression through compositional kernel search. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 1166–1174.
- Ferris, B., Haehnel, D., and Fox, D. (2006). Gaussian processes for signal strength-based location estimation. In *Proceedings of the Conference on Robotics Science and Systems*. Citeseer.
- Gelbart, M. A., Snoek, J., and Adams, R. P. (2014). Bayesian optimization with unknown constraints. *arXiv preprint arXiv:1403.5607*.
- Goodman, N. D .and Mansinghka, V. K., Roy, D., Bonawitz, K., and Tenenbaum, J. (2008). Church: A language for generative models. In *Proceedings of the Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 220–229.
- Kemmler, M., Rodner, E., Wacker, E., and Denzler, J. (2013). One-class classification with gaussian processes. *Pattern Recognition*, 46(12):3507–3518.
- Kennedy, M. C. and O'Hagan, A. (2001). Bayesian calibration of computer models. *Journal of the Royal Statistical Society. Series B, Statistical Methodology*, pages 425–464.
- Kuss, M. and Rasmussen, C. E. (2005). Assessing approximate inference for binary gaussian process classification. *The Journal of Machine Learning Research*, 6:1679–1704.
- Kwan, J., Bhattacharya, S., Heitmann, K., and Habib, S. (2013). Cosmic emulation: The concentration-mass relation for wcdm universes. *The Astrophysical Journal*, 768(2):123.
- Lloyd, J. R., Duvenaud, D., Grosse, R., Tenenbaum, J., and Ghahramani, Z. (2014). Automatic construction and natural-language description of nonparametric regression models. In *Proceedings of the Conference on Artificial Intelligence (AAAI)*.
- Mansinghka, V. K., Selsam, D., and Perov, Y. (2014). Venture: a higher-order probabilistic programming platform with programmable inference. *arXiv preprint arXiv:1404.0099*.

- 1188 McAllester, D., Milch, B., and Goodman, N. D. (2008). Random-world semantics and syntactic  
1189 independence for expressive languages. Technical report.  
1190
- 1191 Neal, R. M. (1995). *Bayesian Learning for Neural Networks*. PhD thesis, University of Toronto.  
1192
- 1193 Neal, R. M. (1997). Monte carlo implementation of gaussian process models for bayesian regression  
and classification. *arXiv preprint physics/9701026*.
- 1194 Norvig, P. (1992). *Paradigms of artificial intelligence programming: case studies in Common LISP*.  
1195 Morgan Kaufmann.
- 1196 Poole, D. (1993). Probabilistic horn abduction and bayesian networks. *Artificial Intelligence*,  
1197 64(1):81–129.
- 1198 Rasmussen, C. E. and Williams, C. K. I. (2006). *Gaussian Processes for Machine Learning (Adap-  
1199 tive Computation and Machine Learning)*. The MIT Press.
- 1200 Sato, T. (1995). A statistical learning method for logic programs with distribution semantics. In *In  
1201 Proceedings of the International Conference on Logic Programming*. Citeseer.
- 1202 Schneider, M. D., Knox, L., Habib, S., Heitmann, K., Higdon, D., and Nakhleh, C. (2008). Simula-  
1203 tions and cosmological inference: A statistical model for power spectra means and covariances.  
1204 *Physical Review D*, 78(6):063529.
- 1205 Snoek, J., Larochelle, H., and Adams, R. P. (2012). Practical bayesian optimization of machine  
1206 learning algorithms. In *Advances in Neural Information Processing Systems (NIPS)*, pages 2951–  
1207 2959.
- 1208 Thompson, W. R. (1933). On the likelihood that one unknown probability exceeds another in view  
1209 of the evidence of two samples. *Biometrika*, pages 285–294.
- 1210 Williams, C. K. I. and Barber, D. (1998). Bayesian classification with gaussian processes. *IEEE  
1211 Transactions on Pattern Analysis and Machine Intelligence*, 20(12):1342–1351.
- 1212
- 1213
- 1214
- 1215
- 1216
- 1217
- 1218
- 1219
- 1220
- 1221
- 1222
- 1223
- 1224
- 1225
- 1226
- 1227
- 1228
- 1229
- 1230
- 1231
- 1232
- 1233
- 1234
- 1235
- 1236
- 1237
- 1238
- 1239
- 1240
- 1241